

# Big Data Analytics

Expand the capacity to implement data science research, data science training, and industry collaborations using cutting-edge computational tools.

Prepared by: TENG CHANTO

---



# 1. What is big data?

- Lots of Data (Terabytes or Petabytes)
- Big data is the term for a collection of data sets so **large and complex** that it becomes **difficult** to process using on-hand database management tools or traditional data processing applications
- The challenges include **capture, curation, storage, search, sharing, transfer, analysis, and visualization**





# Types of Big Data


Structured, Unstructured and Semi-Structured

Semi-Structured Data

Structured Data




Unstructured Data



# IBM Definition of big data

IBM's Definition – Big Data Characteristics


<http://www-01.ibm.com/software/data/bigdata/>



VOLUME



VELOCITY



VARIETY

Min	Max	Mean	SD
4.3	7.9	5.84	0.83
2.0	4.4	3.05	0.43
0.1	2.5	1.20	0.76

VERACITY

# Examples of Big Data

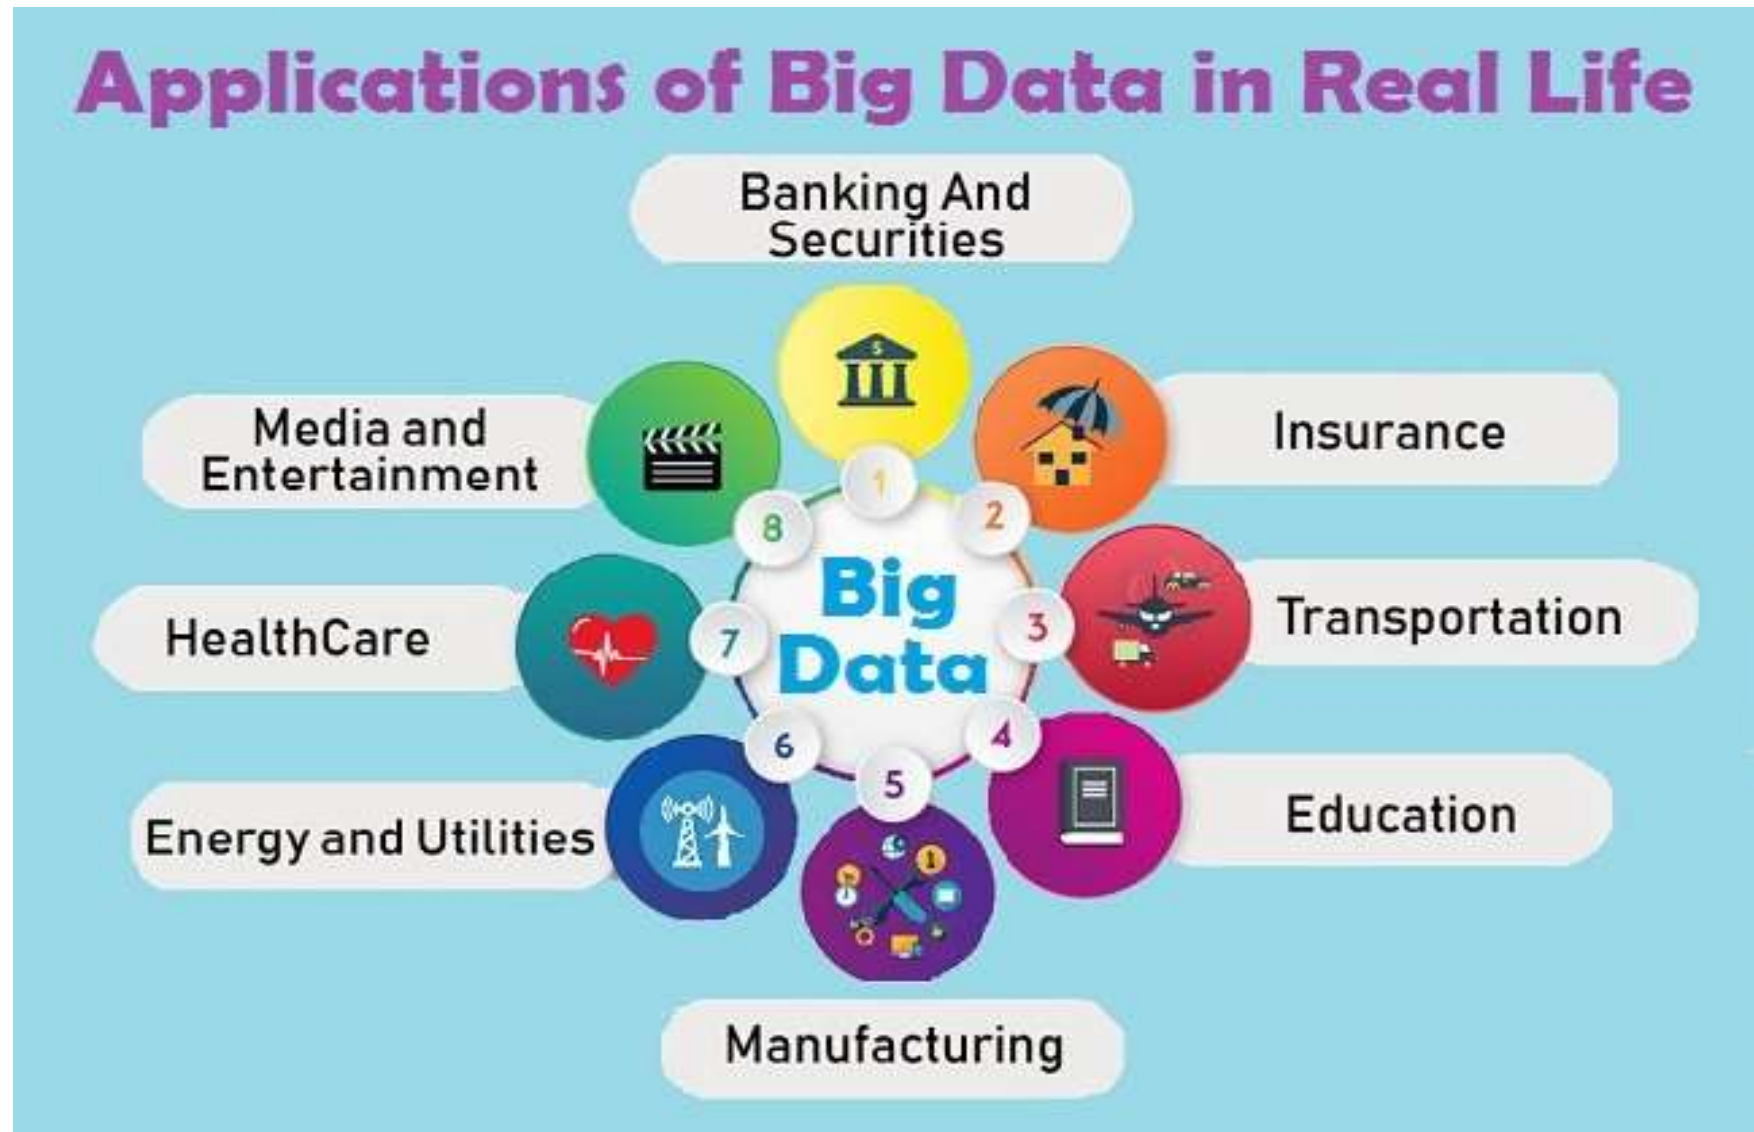
- Walmart handles more than **1 million** customer transactions every hour.
- Facebook stores, accesses, and analyzes **30+ Petabytes** of user generated data.
- **230+ millions** of tweets are created every day.
- More than **5 billion** people are calling, texting, tweeting and browsing on mobile phones worldwide.
- YouTube users upload **48 hours** of new video every minute of the day.
- Amazon handles **15 million** customer click stream user data per day to recommend products.
- **294 billion** emails are sent every day. Services analyses this data to find the spams.
- Modern cars have close to **100 sensors** which monitors fuel level, tire pressure etc. , each vehicle generates a lot of sensor data.

# Challenges with Big Data

- Data Quality
- Discovery
- Storage
- Analytics
- Security
- Lack of Talent



# Big Data Applications



# Big Data in Government

- Cyber security & Intelligence
- Crime Prediction and Prevention
- Pharmaceutical Drug Evaluation
- Scientific Research
- Weather Forecasting
- Tax Compliance
- Traffic Optimization






# Real time scenarios

- Facebook: Ad targeting
- Air traffic
- Electricity grid
- Web & e tailing: recommendations
- Twitter
- Marketing engine

# Sears use Hadoop for analysis

- Insight into data can provide **Business Advantage**.
- Some key early indicators can mean **Fortunes to Business**.
- **More Precise Analysis** with more data.

Case Study: Sears Holding Corporation



*\*Sears was using traditional systems such as Oracle Exadata, Teradata and SAS etc., to store and process the customer activity and sales data.*


# WHY DFS? Distributed file system

Read 1 TB Data



1 Machine

4 I/O Channels  
Each Channel – 100 MB/s



10 Machine

4 I/O Channels  
Each Channel – 100 MB/s




Read 1 TB Data



1 Machine

4 I/O Channels  
Each Channel – 100 MB/s

43 Minutes



10 Machine

4 I/O Channels  
Each Channel – 100 MB/s




Read 1 TB Data



1 Machine

4 I/O Channels  
Each Channel – 100 MB/s

43 Minutes




10 Machine

4 I/O Channels  
Each Channel – 100 MB/s

4.3 Minutes

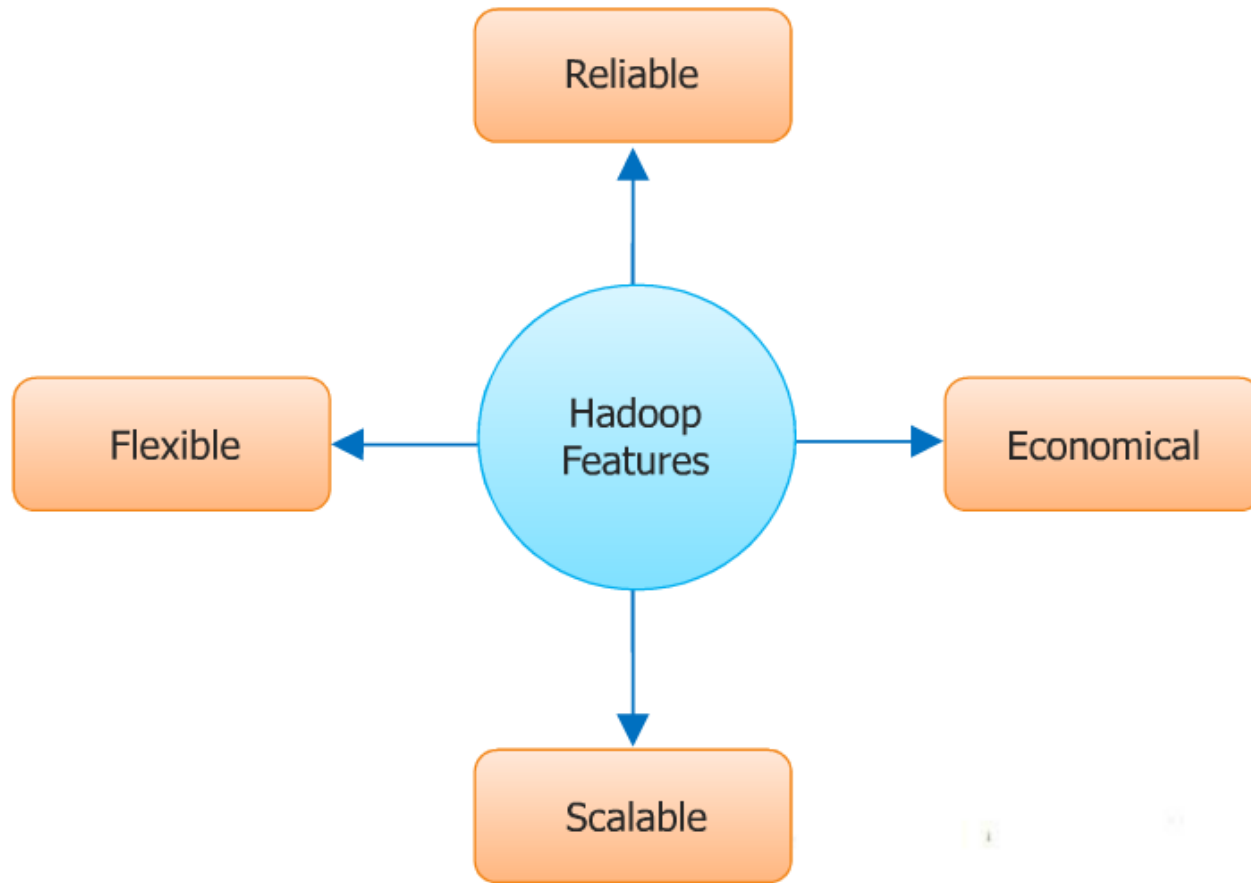
# What is Hadoop?

- Apache Hadoop is a [framework](#) that allows for the distributed processing of large data sets across clusters of commodity computers using a simple programming model.
- It is an [Open-source Data Management](#) with scale-out storage and distributed processing.





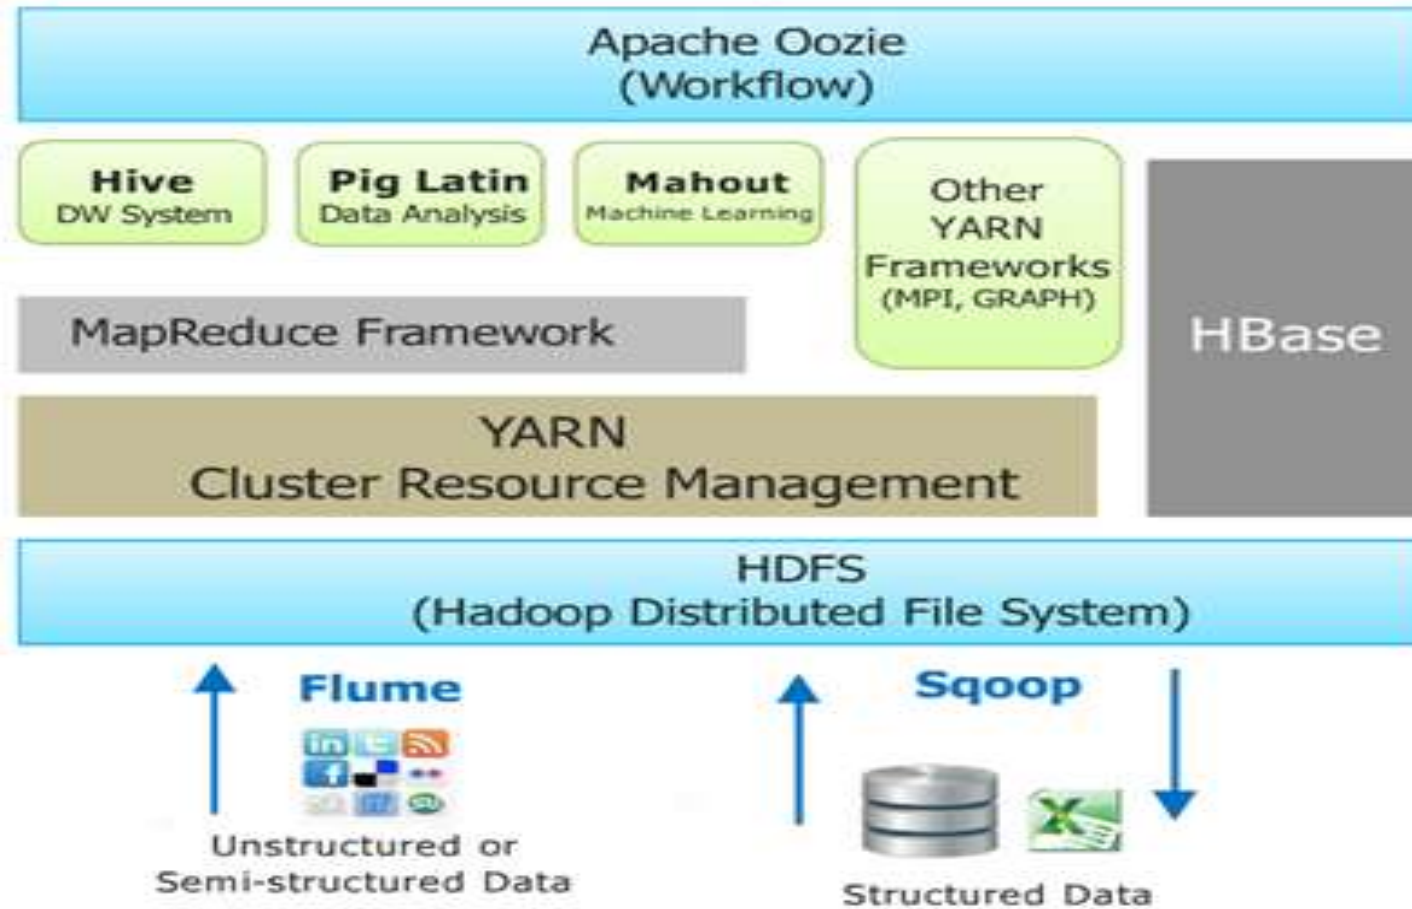
# Characteristics



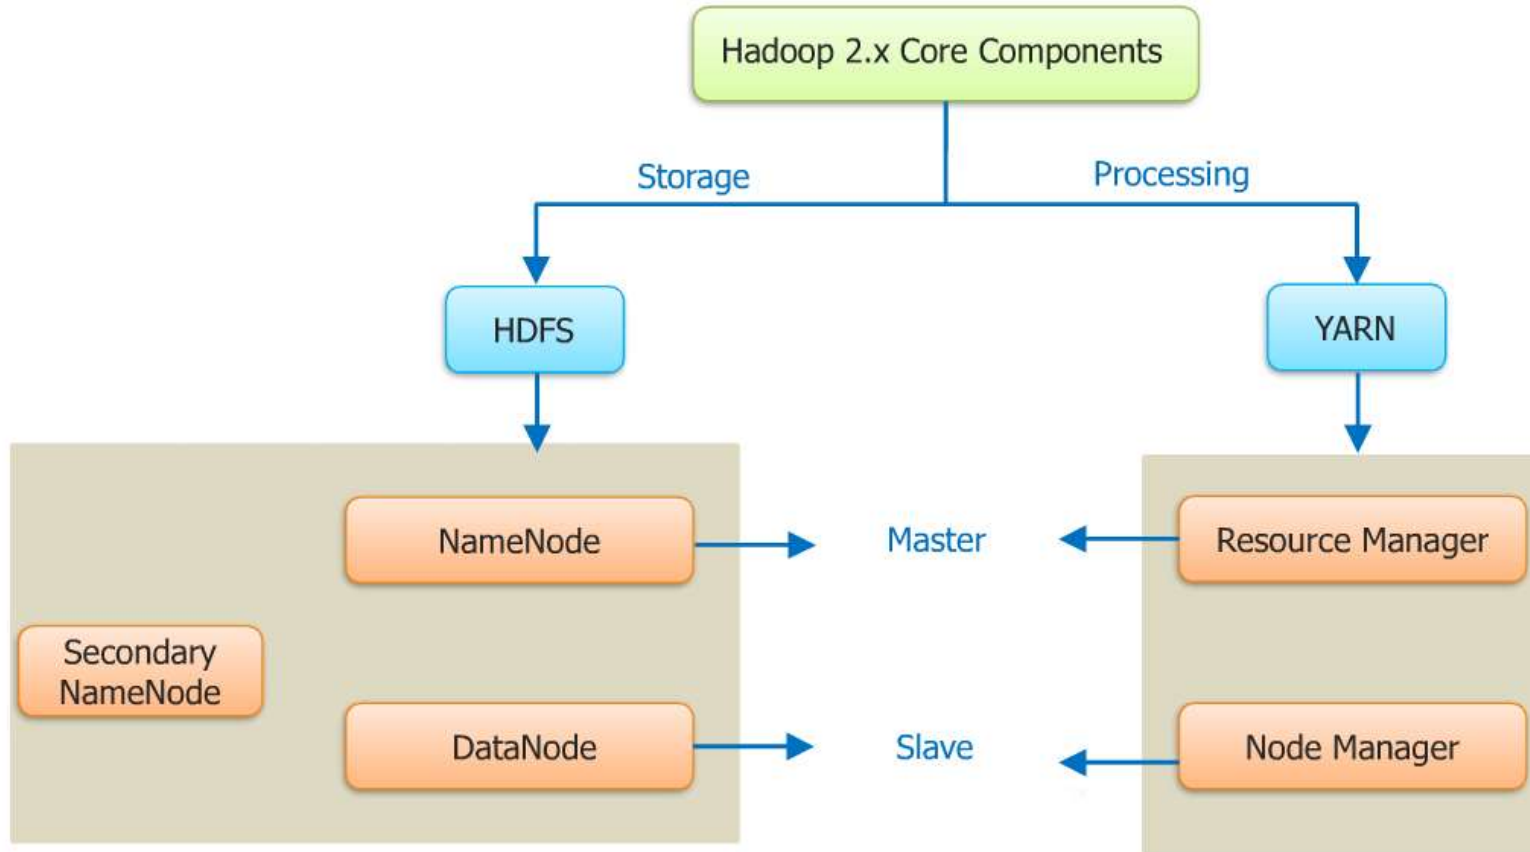
# Hadoop-its about scale and structure

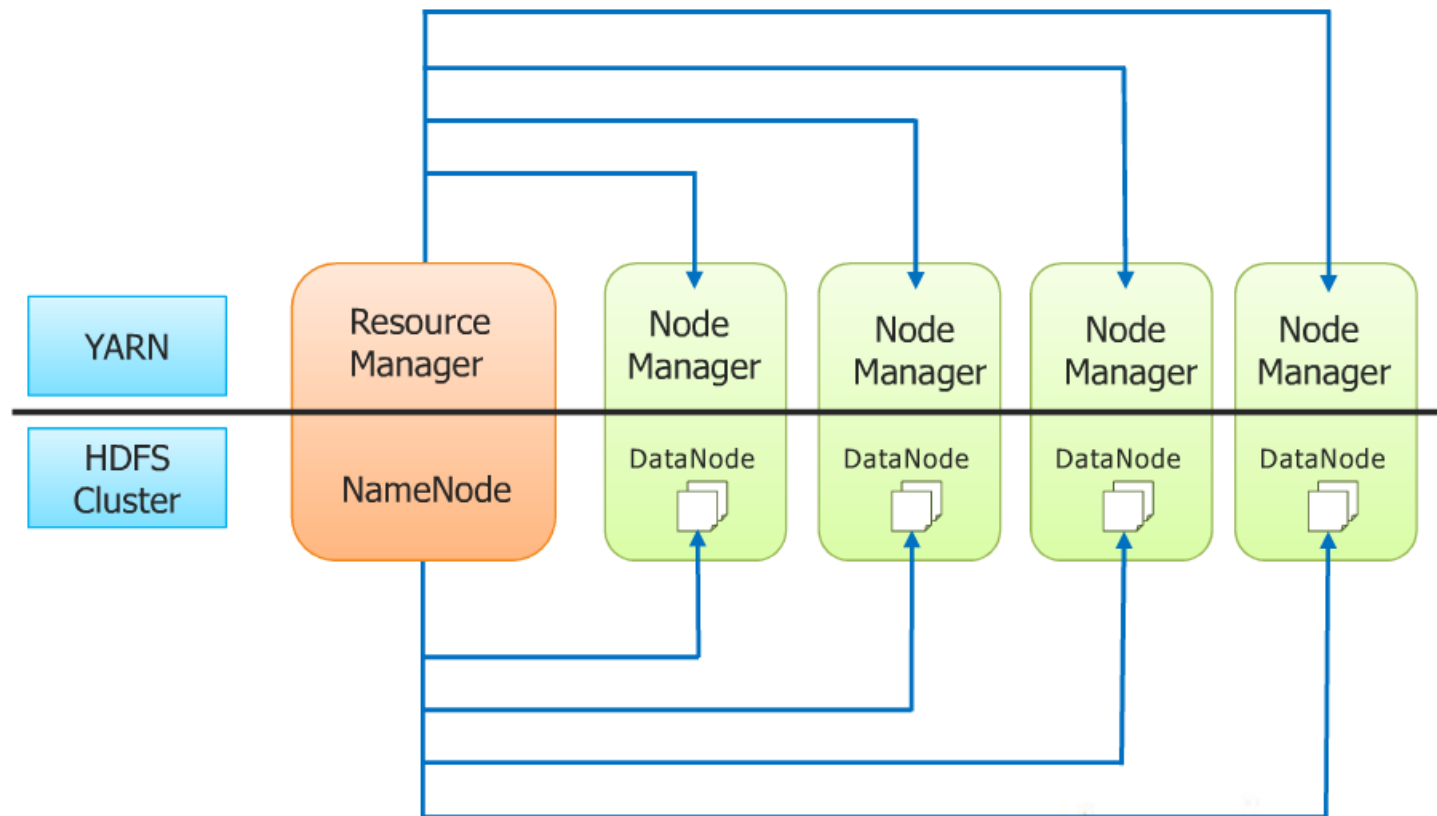
RDBMS ← HADOOP →		
Structured	Data Types	Multi and Unstructured
Limited, No Data Processing	Processing	Processing coupled with Data
Standards & Structured	Governance	Loosely Structured
Required On Write	Schema	Required On Read
Reads are Fast	Speed	Writes are Fast
Software License	Cost	Support Only
Known Entity	Resources	Growing, Complexities, Wide
OLTP Complex ACID Transactions Operational Data Store	Best Fit Use	Data Discovery Processing Unstructured Data Massive Storage/Processing

# Hadoop ecosystem



# Hadoop 2.X core components





# Main components of HDFS

## → NameNode:

- » Master of the system
- » Maintains and manages the blocks which are present on the DataNodes



## → DataNodes:

- » Slaves which are deployed on each machine and provide the actual storage
- » Responsible for serving read and write requests for the clients





# Name Node Metadata

## → Meta-data in Memory

- » The entire metadata is in main memory
- » No demand paging of FS meta-data

## → Types of Metadata

- » List of files
- » List of Blocks for each file
- » List of DataNode for each block
- » File attributes, e.g. access time, replication factor

## → A Transaction Log

- » Records file creations, file deletions etc.

NameNode  
(Stores metadata only)

METADATA:

/user/doug/hinfo -> 1 3 5

/user/doug/pdetail -> 4 2

**NameNode:**



Keeps track of overall file directory structure and the placement of Data Block

# File Blocks

→ By Default, block size is **128mb** in Hadoop 2.x and **64mb** in Hadoop 1.x

## Hadoop 2.x

200mb – abc.txt

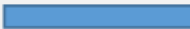
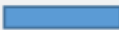
128mb – Block 1   
72mb – Block 2 

300mb – emp.dat





128mb – Block1   
128mb – Block2   
44 mb – Block3 

## Hadoop 1.x

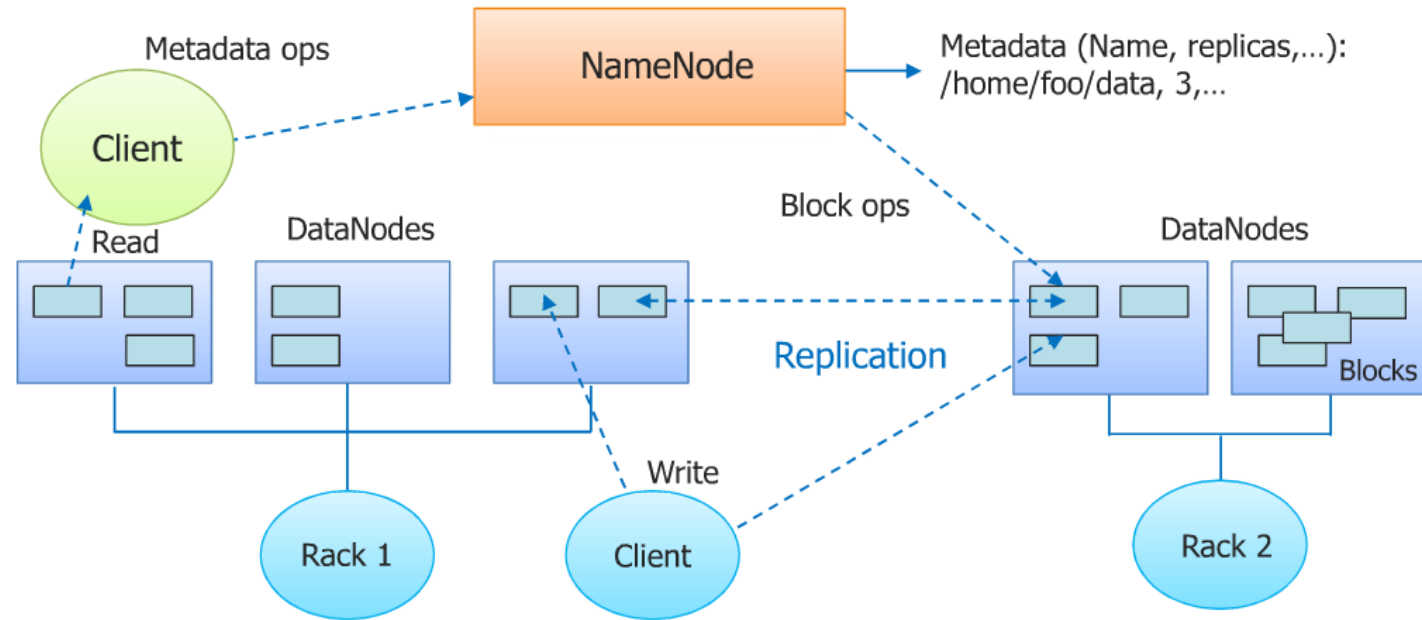
100mb – abc.txt

64mb – Block 1   
36mb – Block 2 

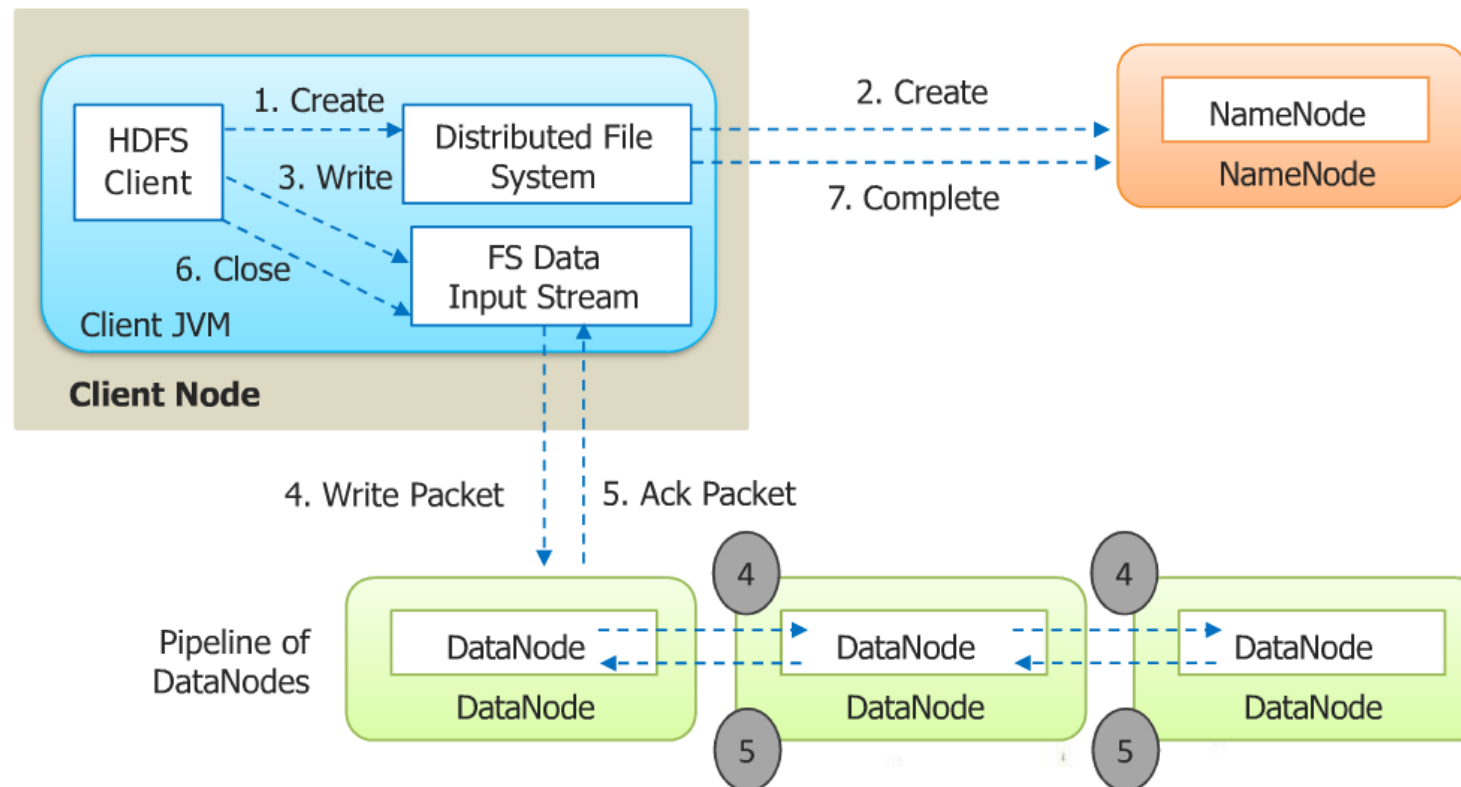
200mb – emp.dat

64mb – Block1   
64mb – Block2   
64mb – Block3   
8mb – Block4 

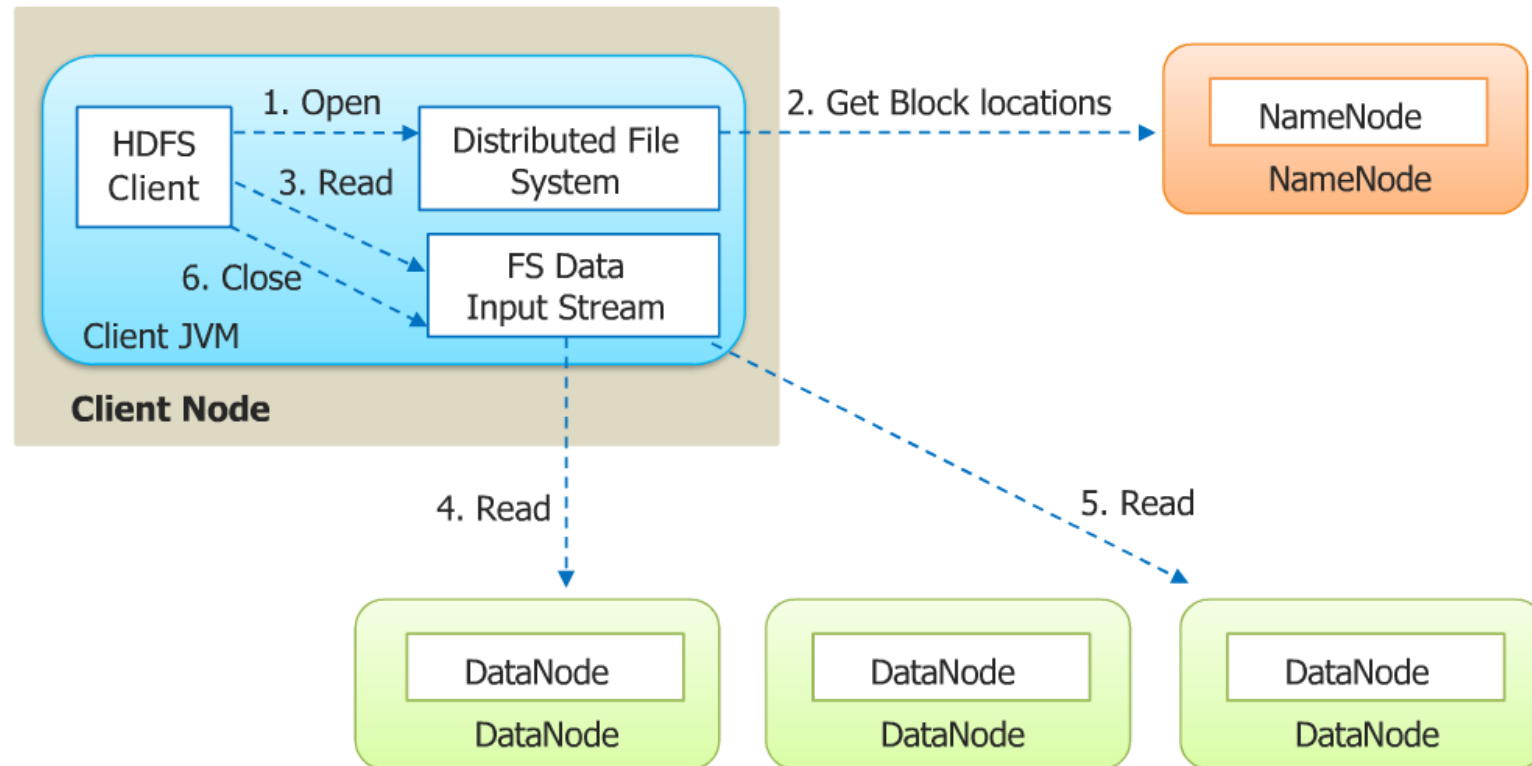
# HDFS Architecture

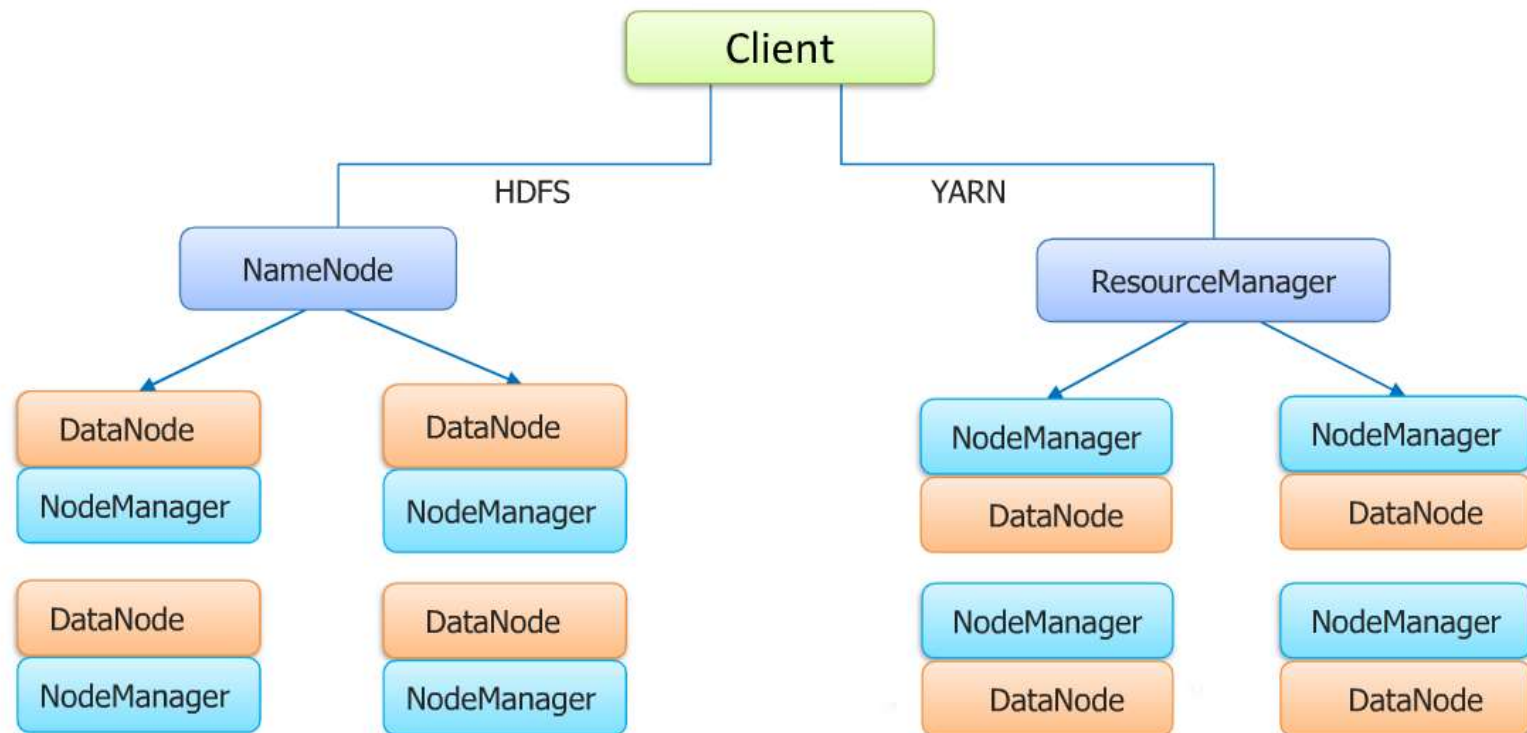


# Anatomy of a file write



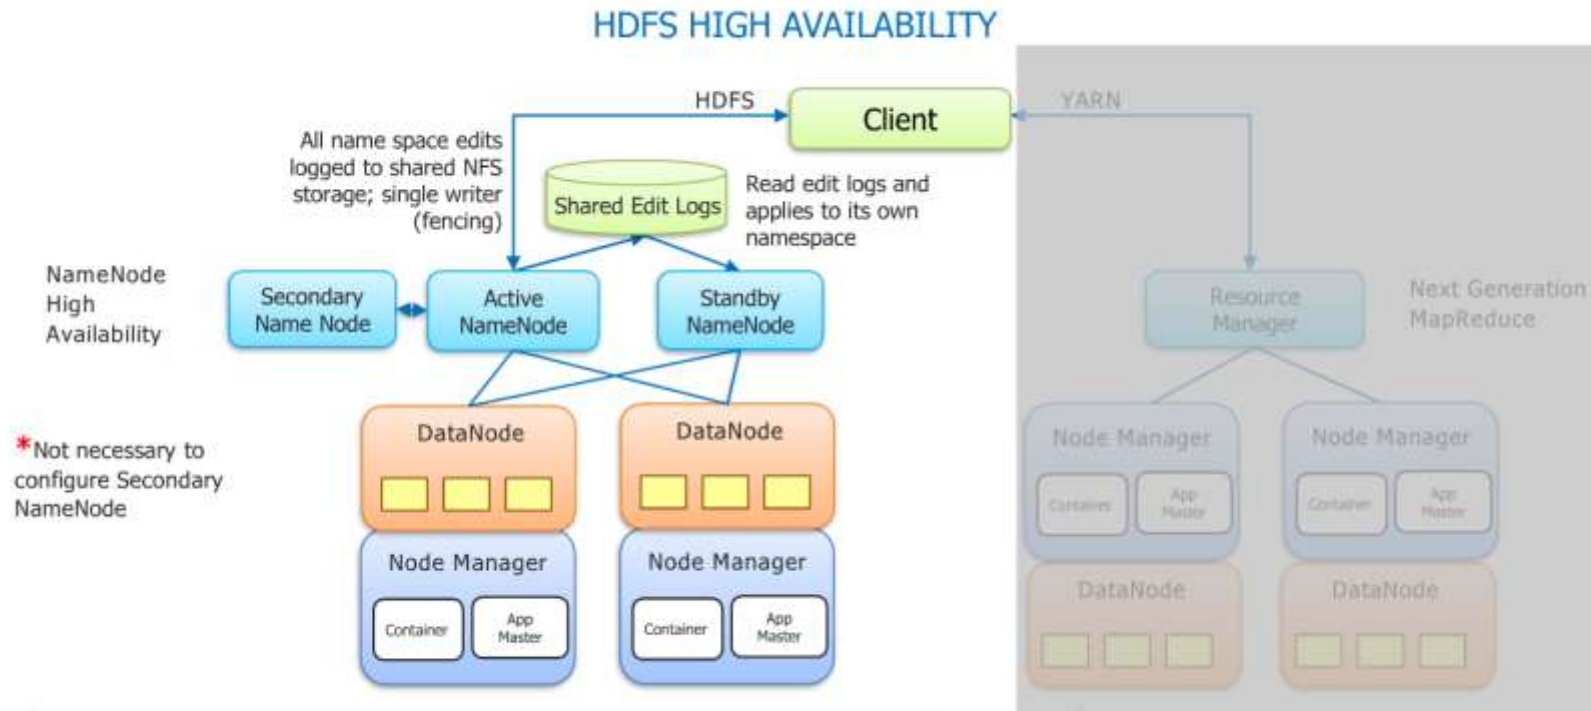
# Anatomy of a file read





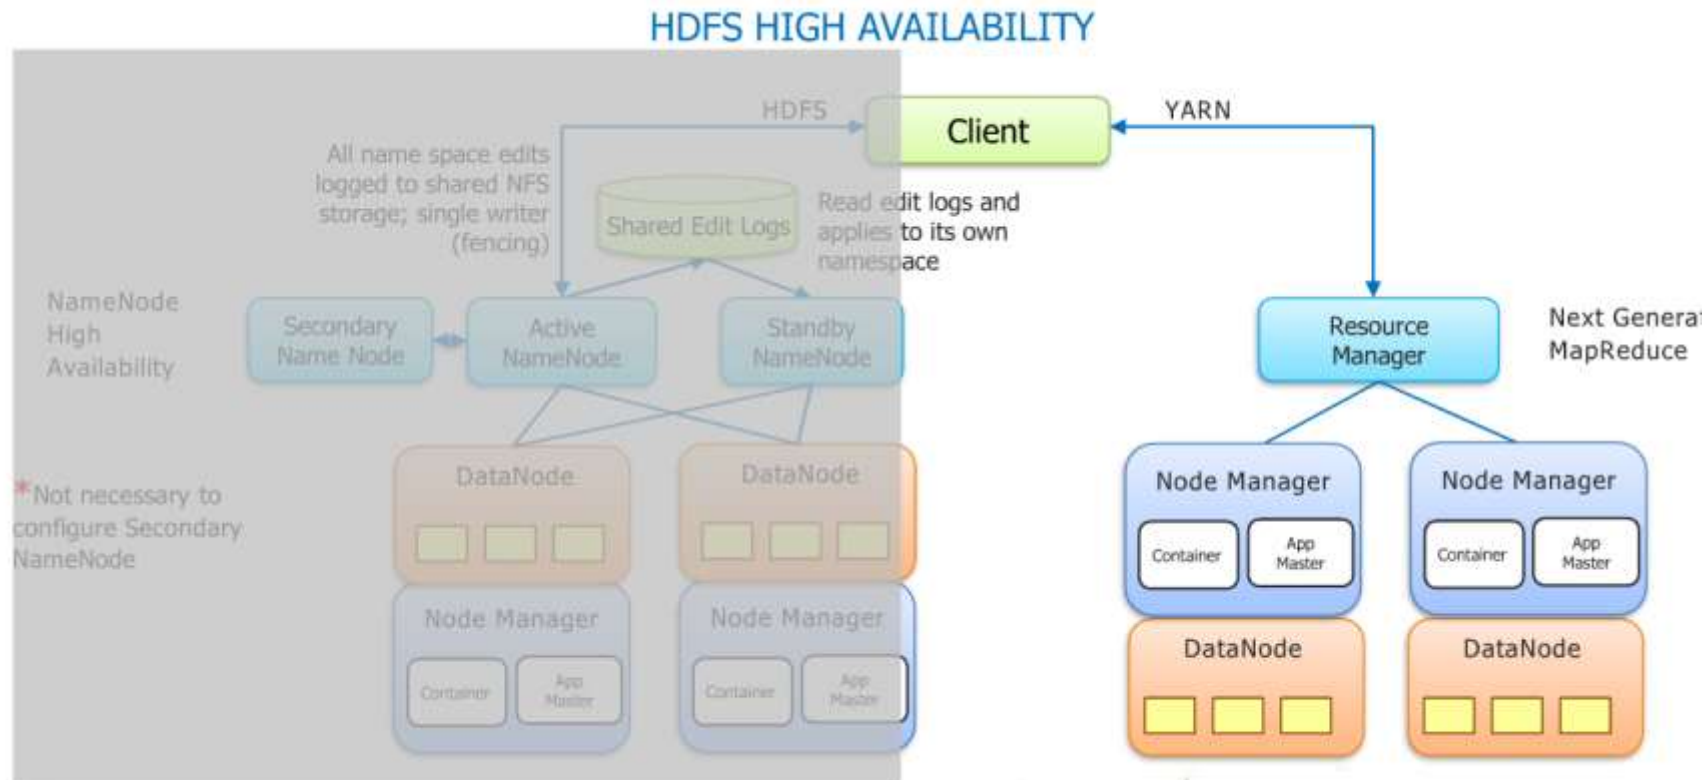


# Hadoop 2.x-High availability

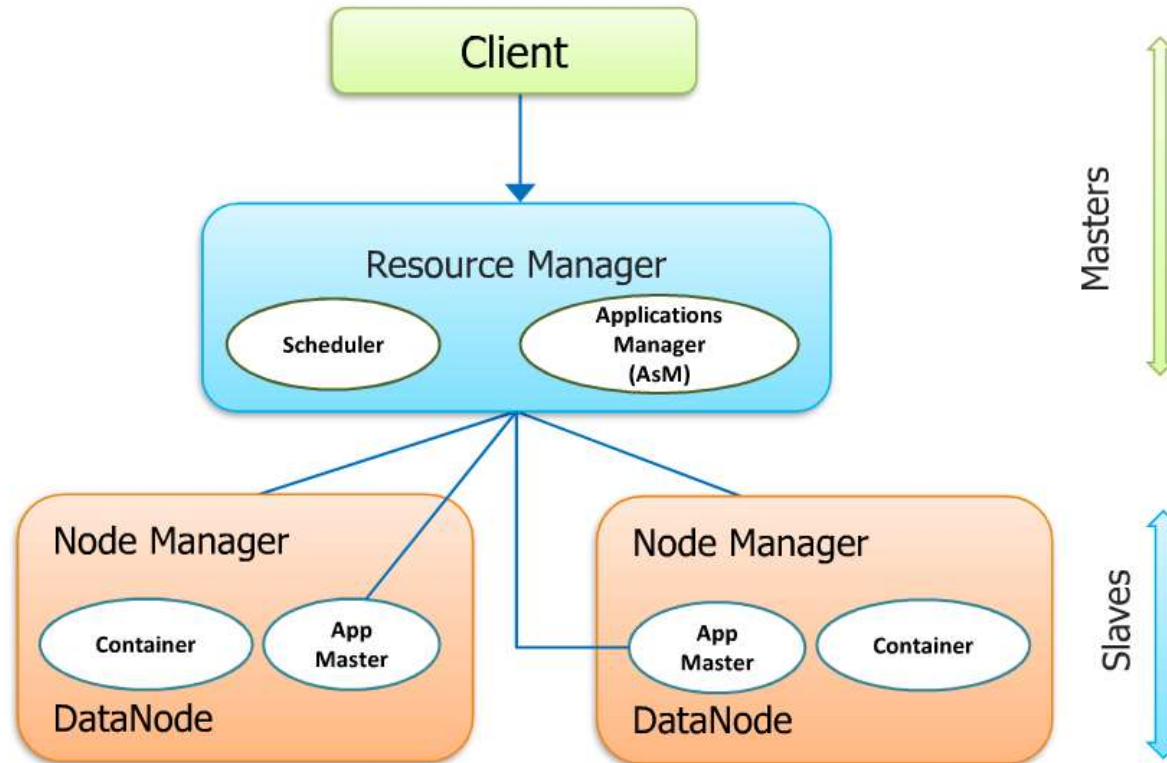


<http://hadoop.apache.org/docs/stable2/hadoop-yarn/hadoop-yarn-site/HDFSHighAvailabilityWithNFS.html>

# Hadoop 2.x-Resource management

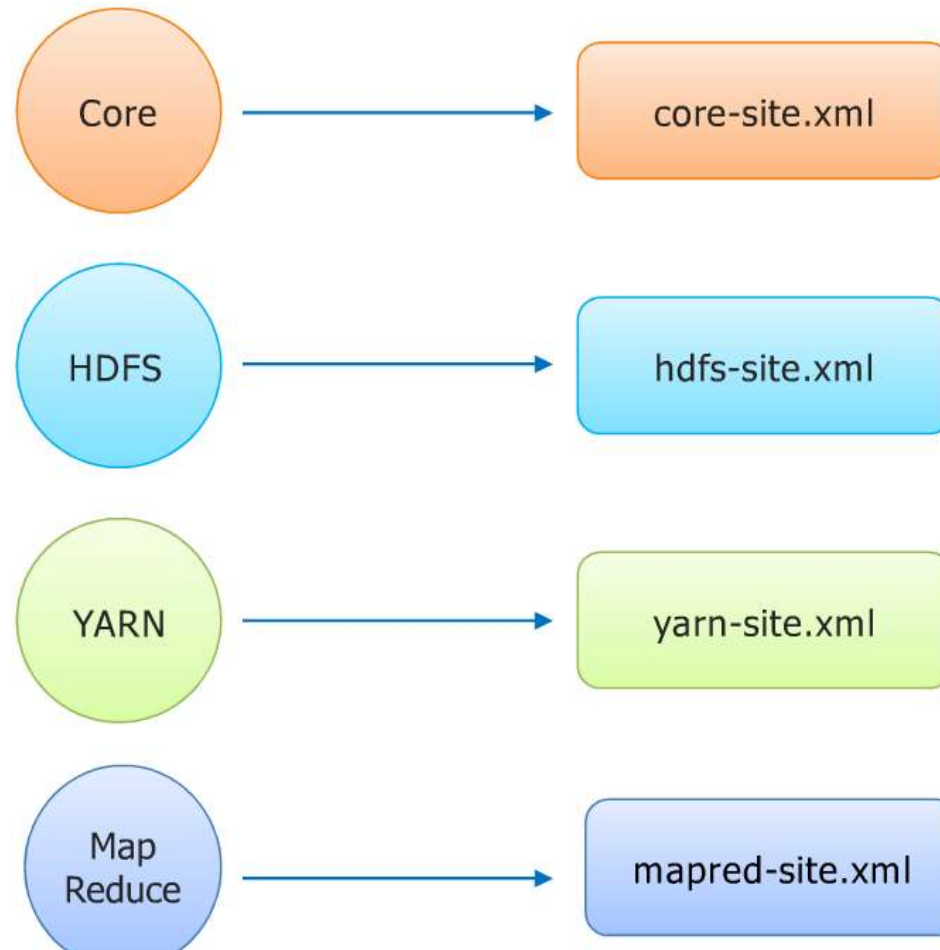


<http://hadoop.apache.org/docs/stable2/hadoop-yarn/hadoop-yarn-site/HDFSHighAvailabilityWithNFS.html>



YARN – Yet Another Resource Negotiator

# Configuration files



# Core-site.xml

-----core-site.xml-----

```
<?xml version="1.0" encoding="UTF-8"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!-- core-site.xml -->
<configuration>
  <property>
    <name>fs.defaultFS</name>
    <value>hdfs://localhost:9000</value>
  </property>
</configuration>
```

The name of the default file system. The url's authority is used to determine the host, port, etc. for a filesystem.

-----core-site.xml-----

# hdfs-site.xml

```
<?xml version="1.0" encoding="UTF-8"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!-- hdfs-site.xml -->
<configuration>
  <property>
    <name>dfs.replication</name>
    <value>1</value>
  </property>
  <property>
    <name>dfs.permissions</name>
    <value>false</value>
  </property>
  <property>
    <name>dfs.namenode.name.dir</name>
    <value>/home/edureka/hadoop-2.2.0/hadoop2_data/hdfs/namenode</value>
  </property>
  <property>
    <name>dfs.datanode.data.dir</name>
    <value>/home/edureka/hadoop-2.2.0/hadoop2_data/hdfs/datanode</value>
  </property>
</configuration>
```

Determines where on the local filesystem the DFS name node should store the name table(fsimage).

If "true", enable permission checking in HDFS. If "false", permission checking is turned off.

Determines where on the local filesystem the DFS name node should store the name table(fsimage).

Determines where on the local filesystem an DFS data node should store its blocks.



# mapred-site.xml

-----mapred-site.xml-----

```
<?xml version="1.0" encoding="UTF-8"?>
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>
<!-- mapred-site.xml -->
<configuration>
  <property>
    <name>mapreduce.framework.name</name>
    <value>yarn</value>
  </property>
</configuration>
```

The runtime framework for  
executing MapReduce jobs.  
Can be one of local, classic  
or yarn.

-----mapred-site.xml-----

# yarn-site.xml

```
-----yarn-site.xml-----  
<?xml version="1.0" encoding="UTF-8"?>  
<?xml-stylesheet type="text/xsl" href="configuration.xsl"?>  
<!-- yarn-site.xml -->  
<configuration>  
  <property>  
    <name>yarn.nodemanager.aux-services</name>  
    <value>mapreduce_shuffle</value>  
  </property>  
  <property>  
    <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>  
    <value>org.apache.hadoop.mapred.ShuffleHandler</value>  
  </property>  
</configuration>
```

The auxiliary service name.

The auxiliary service class to use.

-----yarn-site.xml-----

# Slaves and masters

Two files are used by the startup and shutdown commands:


Slaves

→ Contains a list of hosts, one per line, that are to host DataNode and NodeManager servers.

Masters

→ Contains a list of hosts, one per line, that are to host Secondary NameNode servers.

# Pre process run time environment

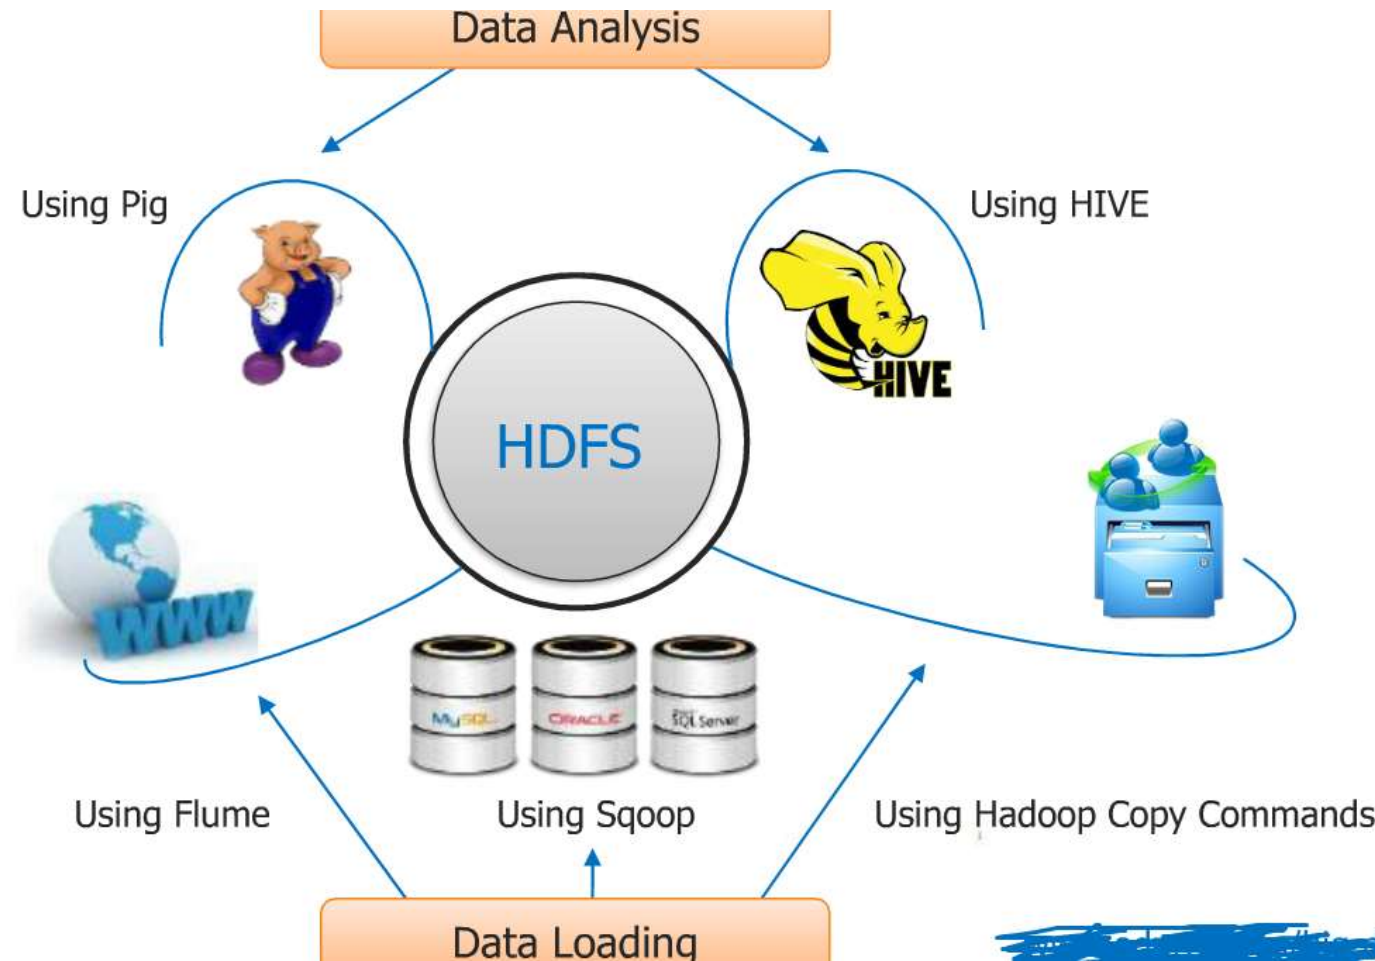


- This file also offers a way to provide custom parameters for each of the servers.
- [Hadoop-env.sh](#) is sourced by all of the Hadoop Core scripts provided in the hadoop directory which is present in hadoop installation directory (`hadoop-2.2.0/etc/hadoop`).
- [Examples of environment variables that you can specify:](#)  

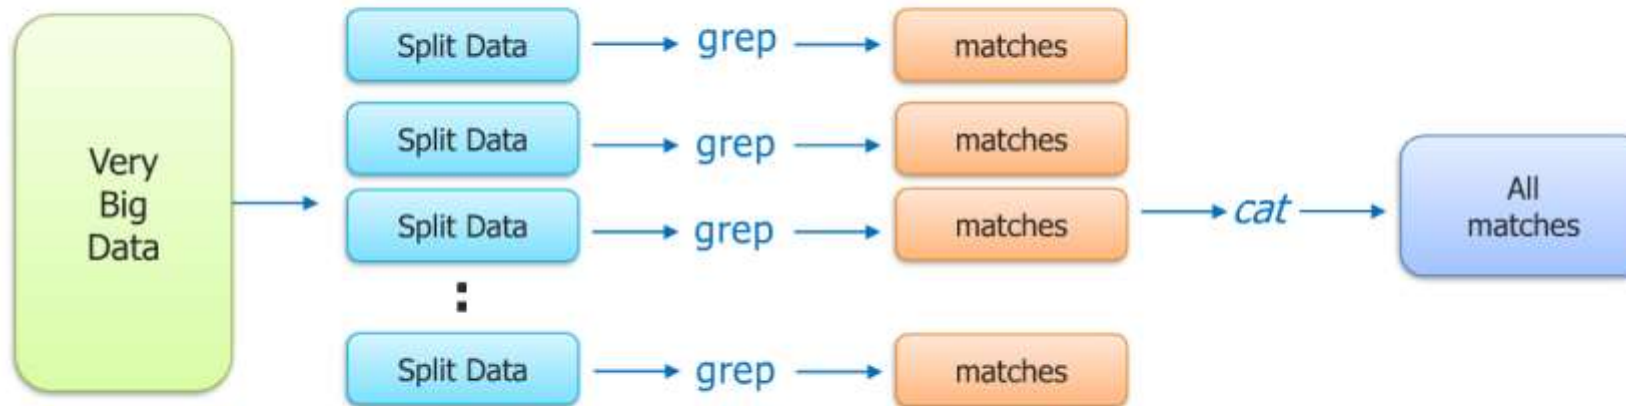
```
export HADOOP_HEAPSIZE="512"
```

```
export HADOOP_DATANODE_HEAPSIZE="128"
```

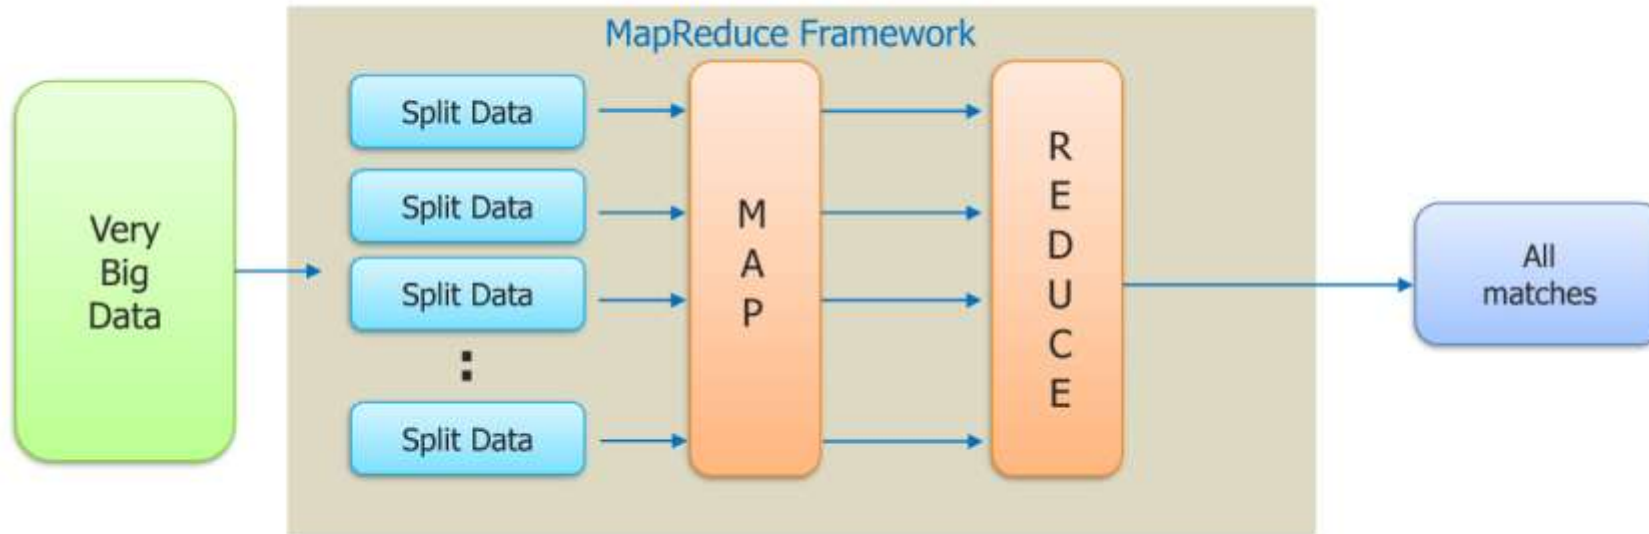
# Data loading techniques and analysis



# The Traditional way



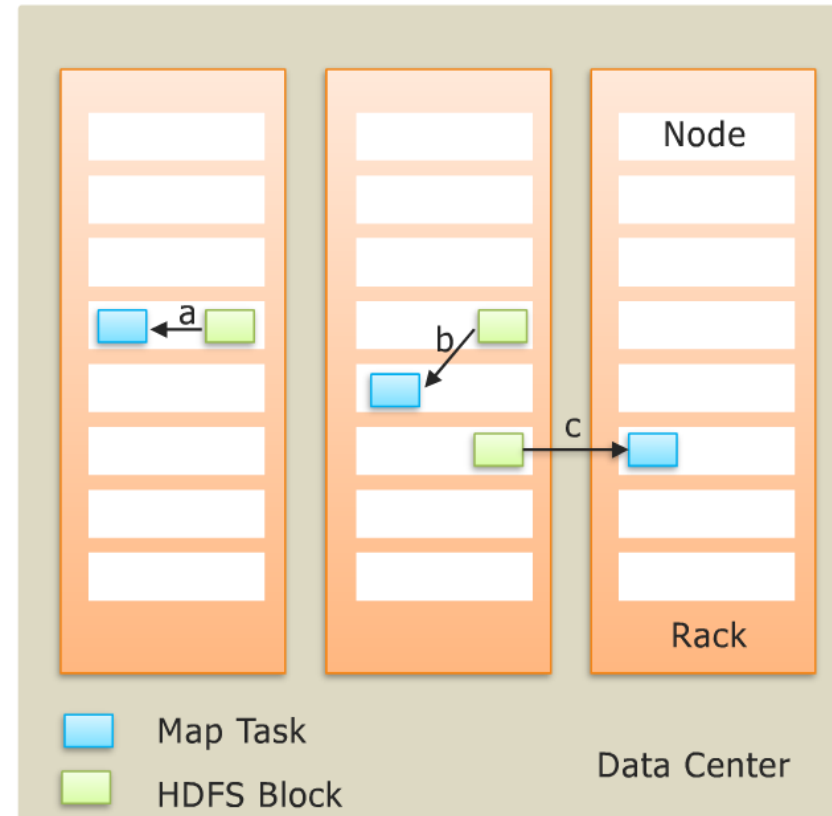
# Map Reduce way



# Why MAP REDUCE

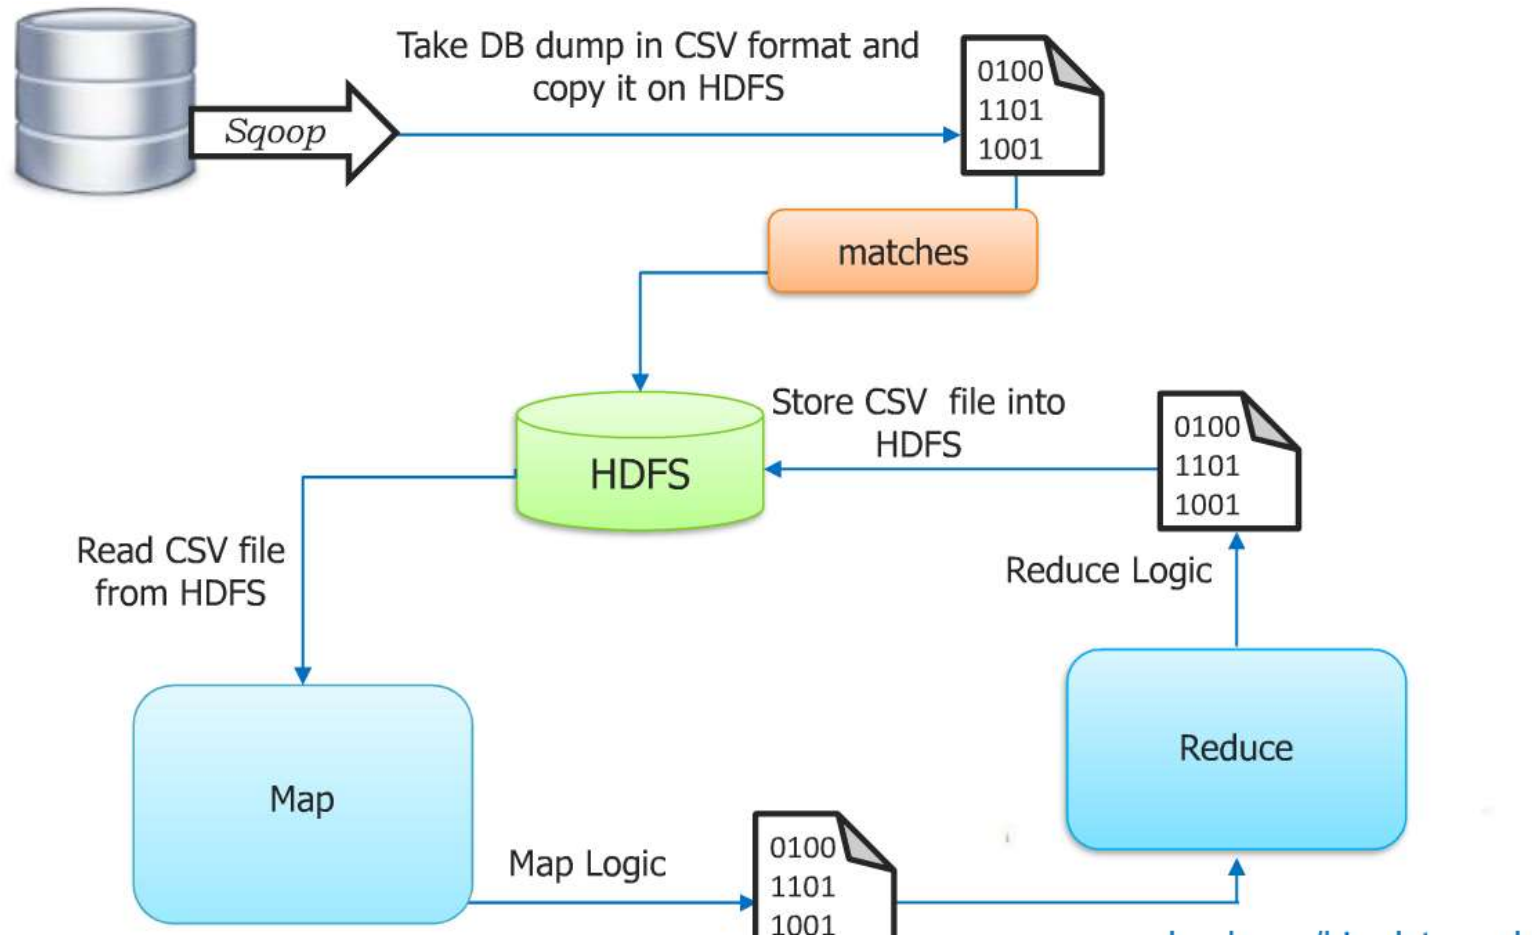
→ Two biggest Advantages:

- » Taking processing to the data
- » Processing data in parallel

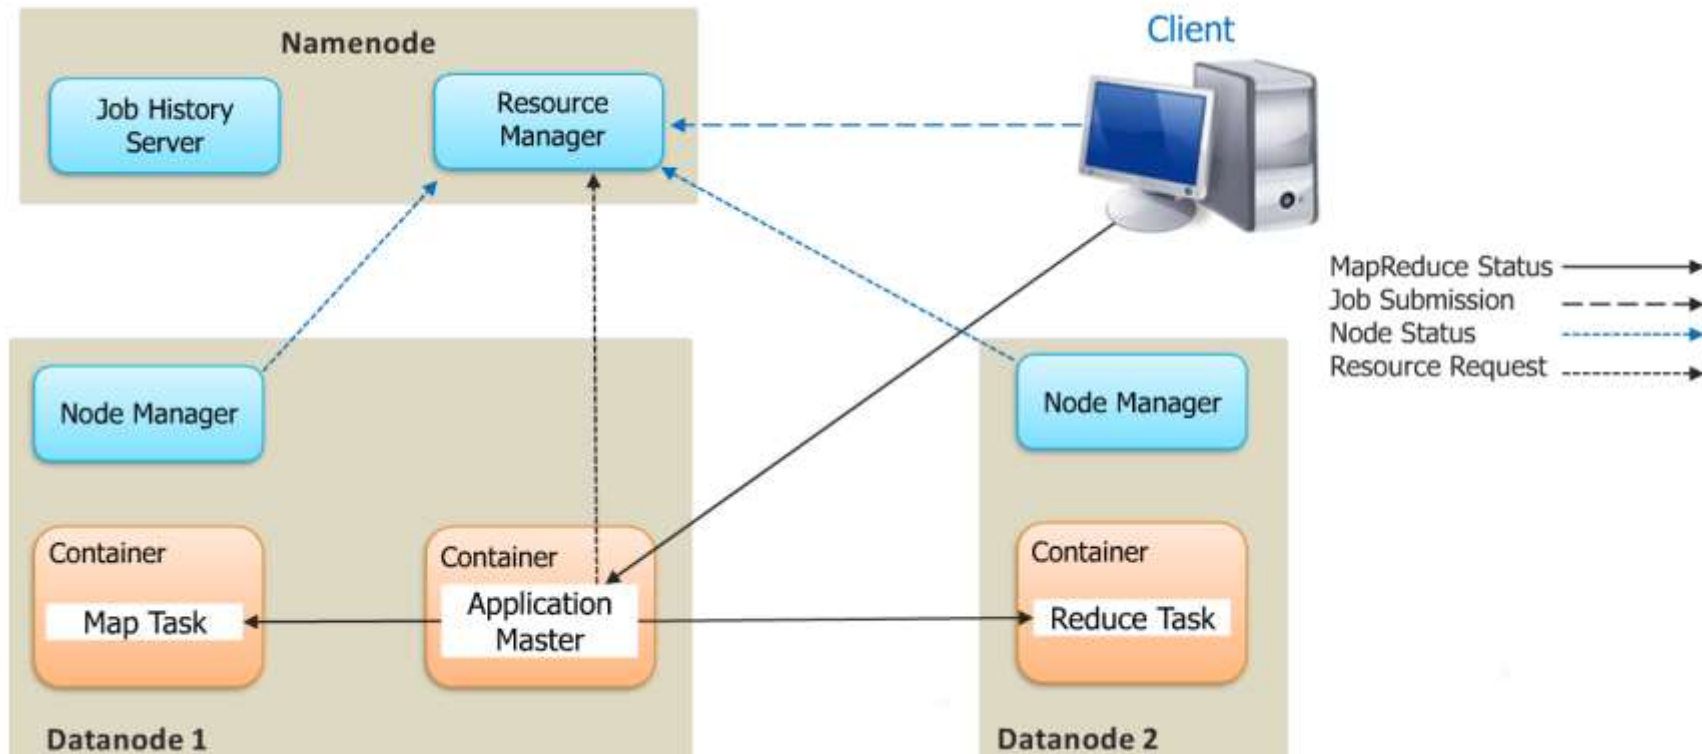


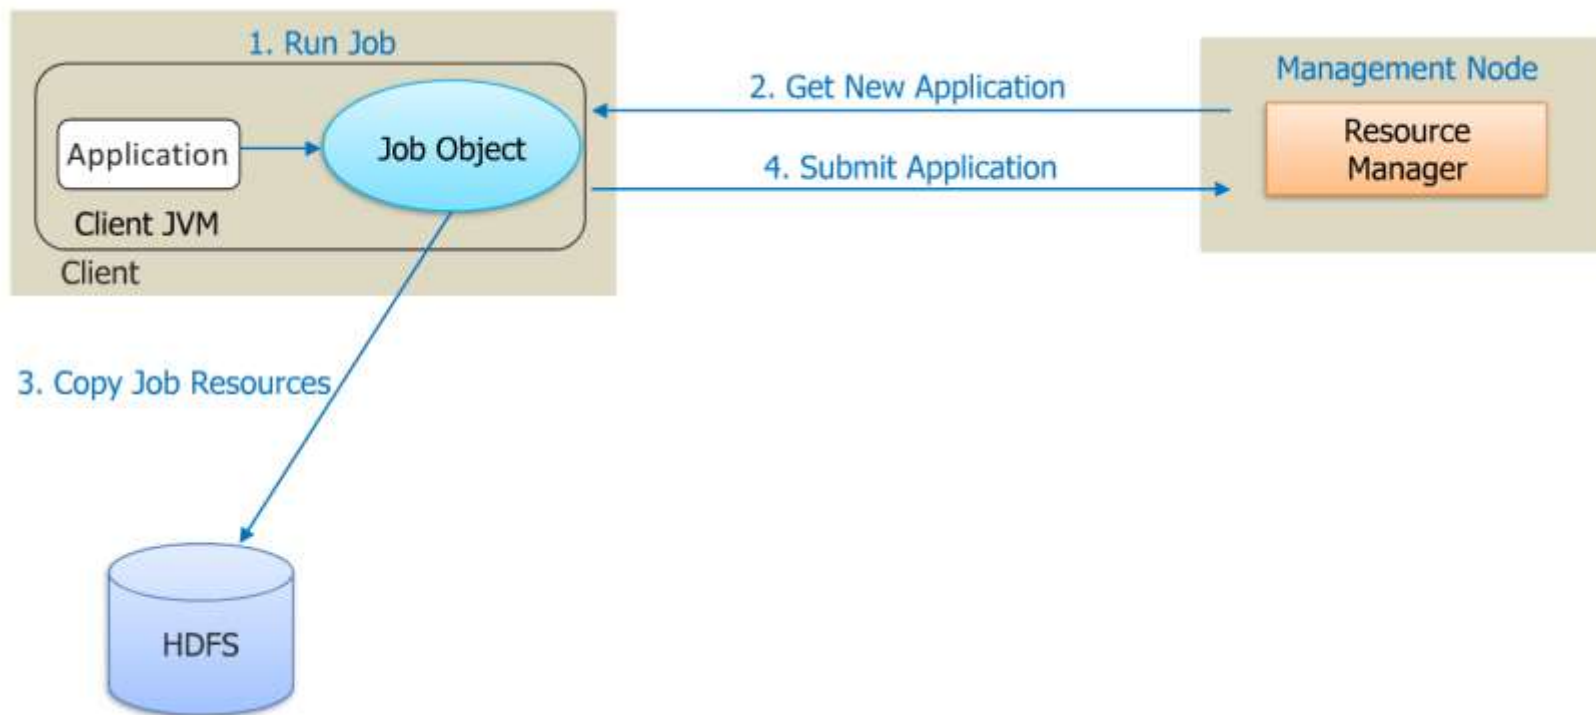


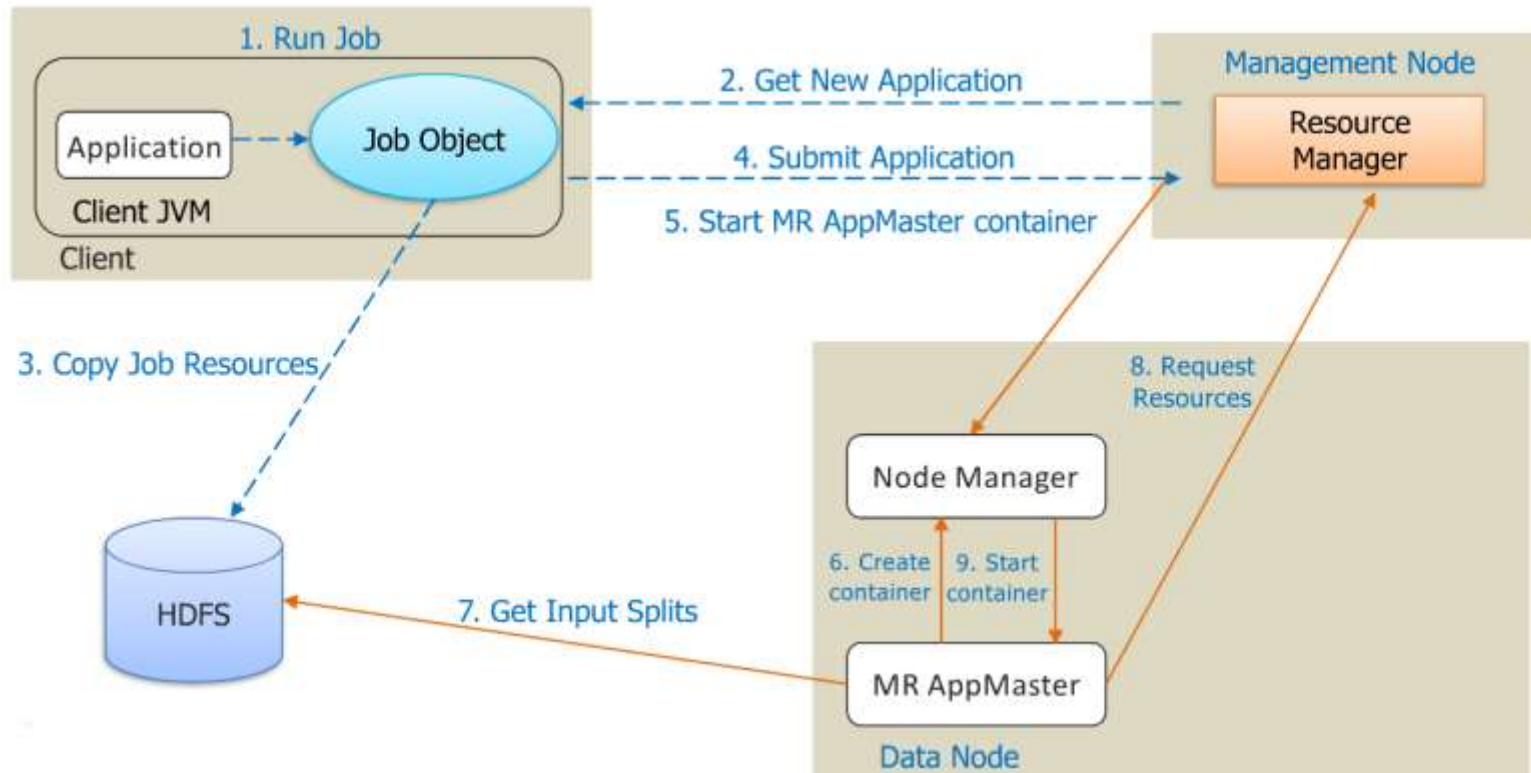
# Solving the problem with map reduce

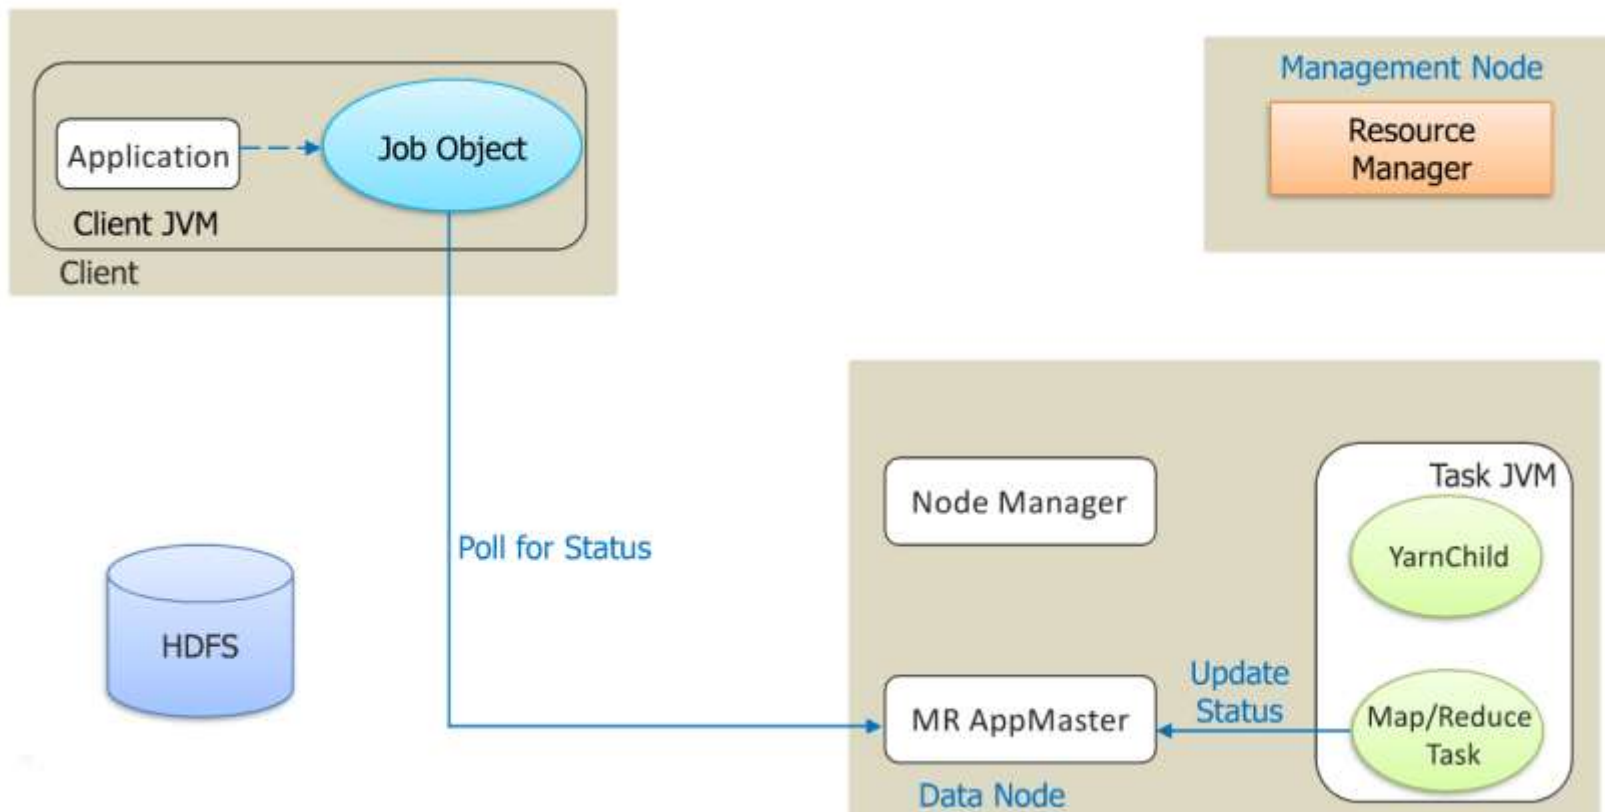


# Hadoop 2.x map reduce architecture



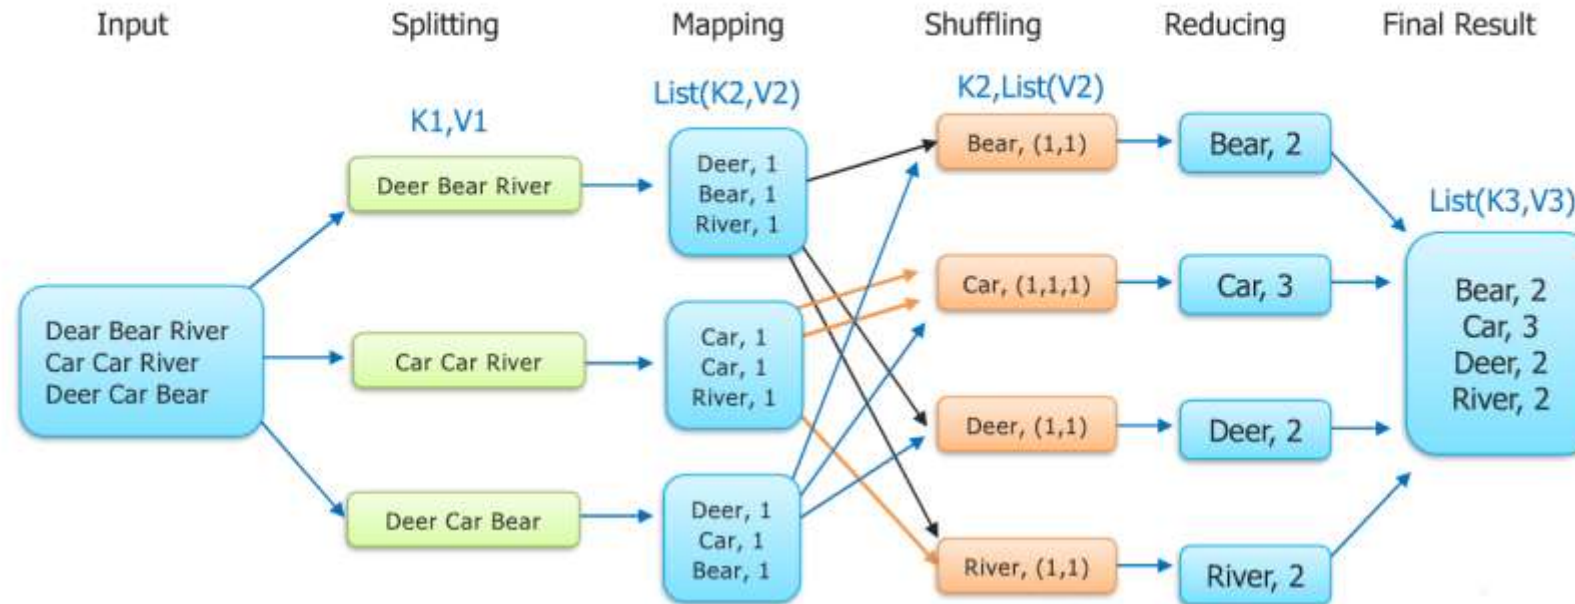







# Map reduce paradigm


The Overall MapReduce Word Count Process



# Anatomy of map reduce program




# Input splits



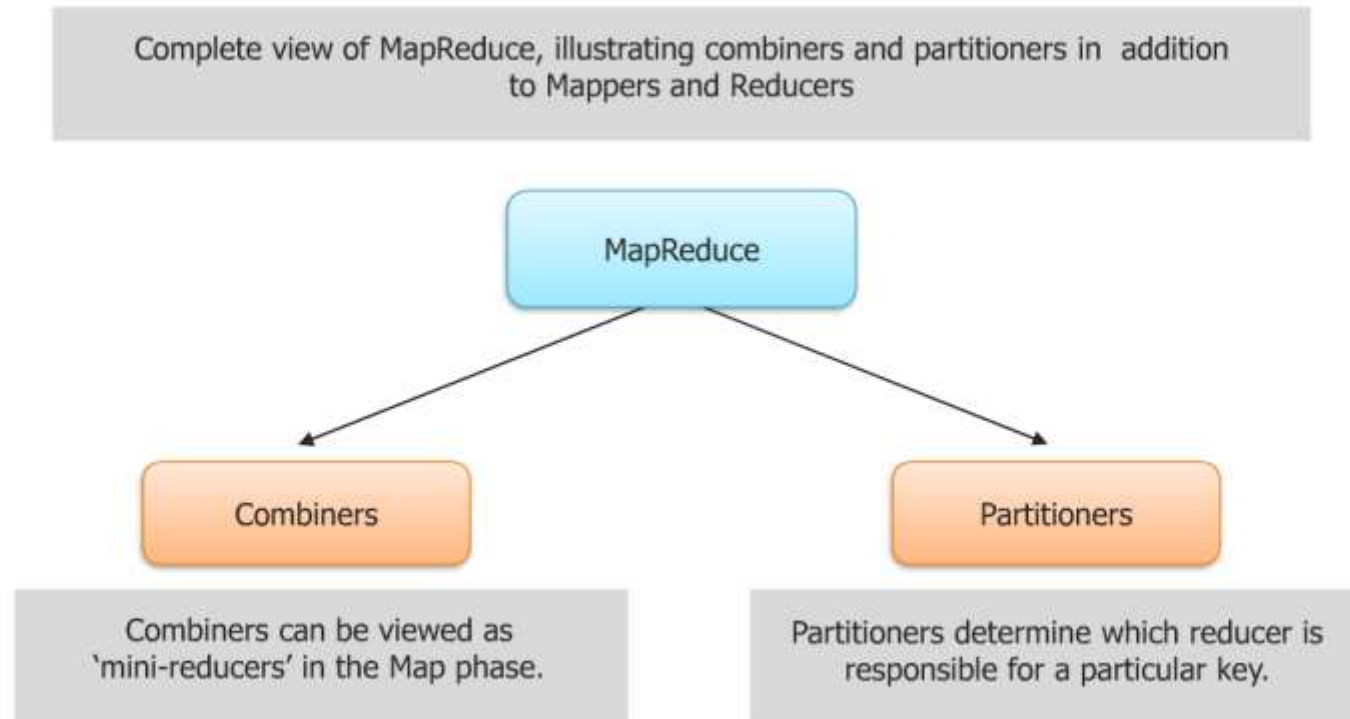


# Map reduce job submission flow

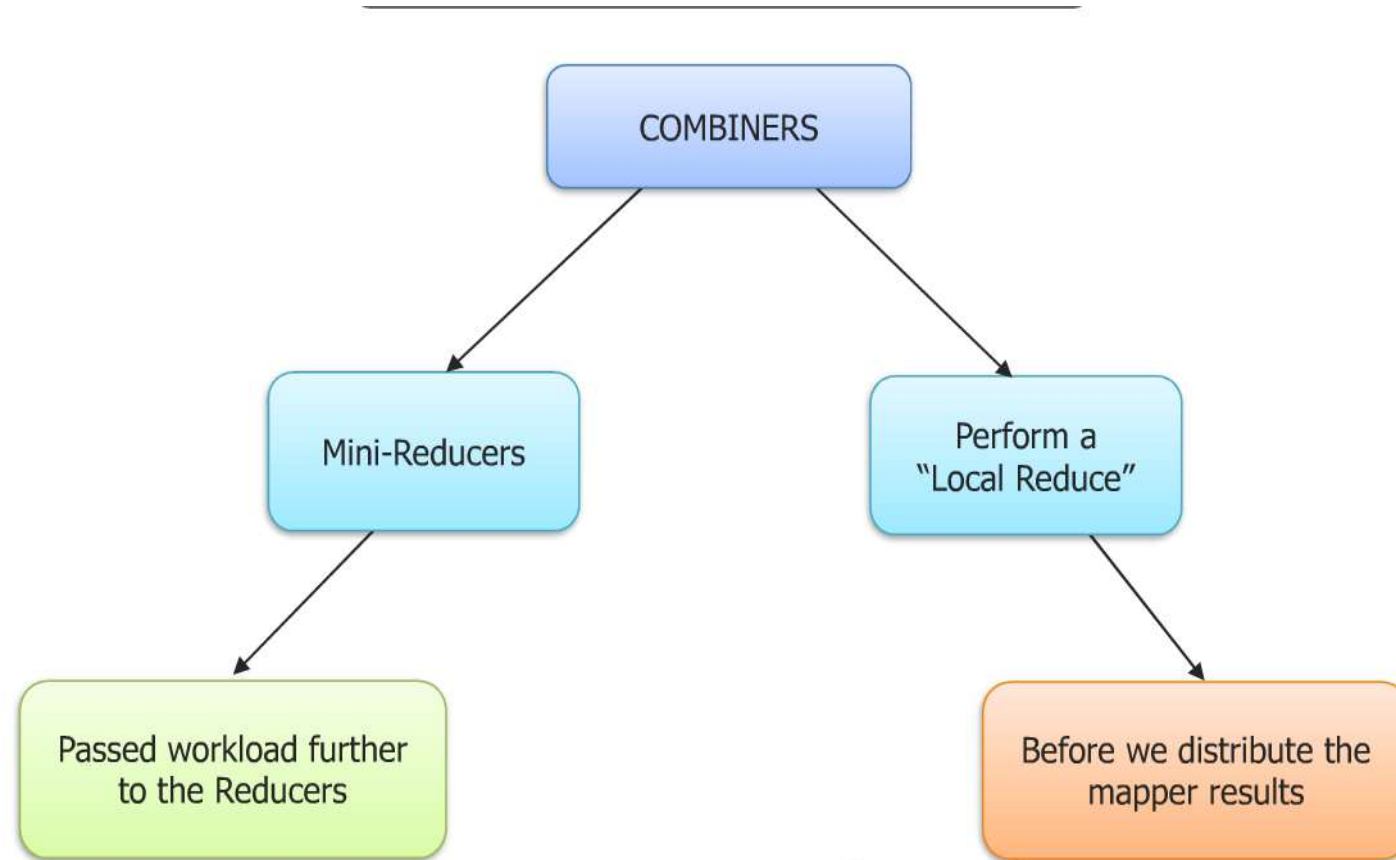
- Input data is distributed to nodes
- Each map task works on a "split" of data
- Mapper outputs intermediate data
- Data exchange between nodes in a "shuffle" process
- Intermediate data of the same key goes to the same reducer
- Reducer output is stored



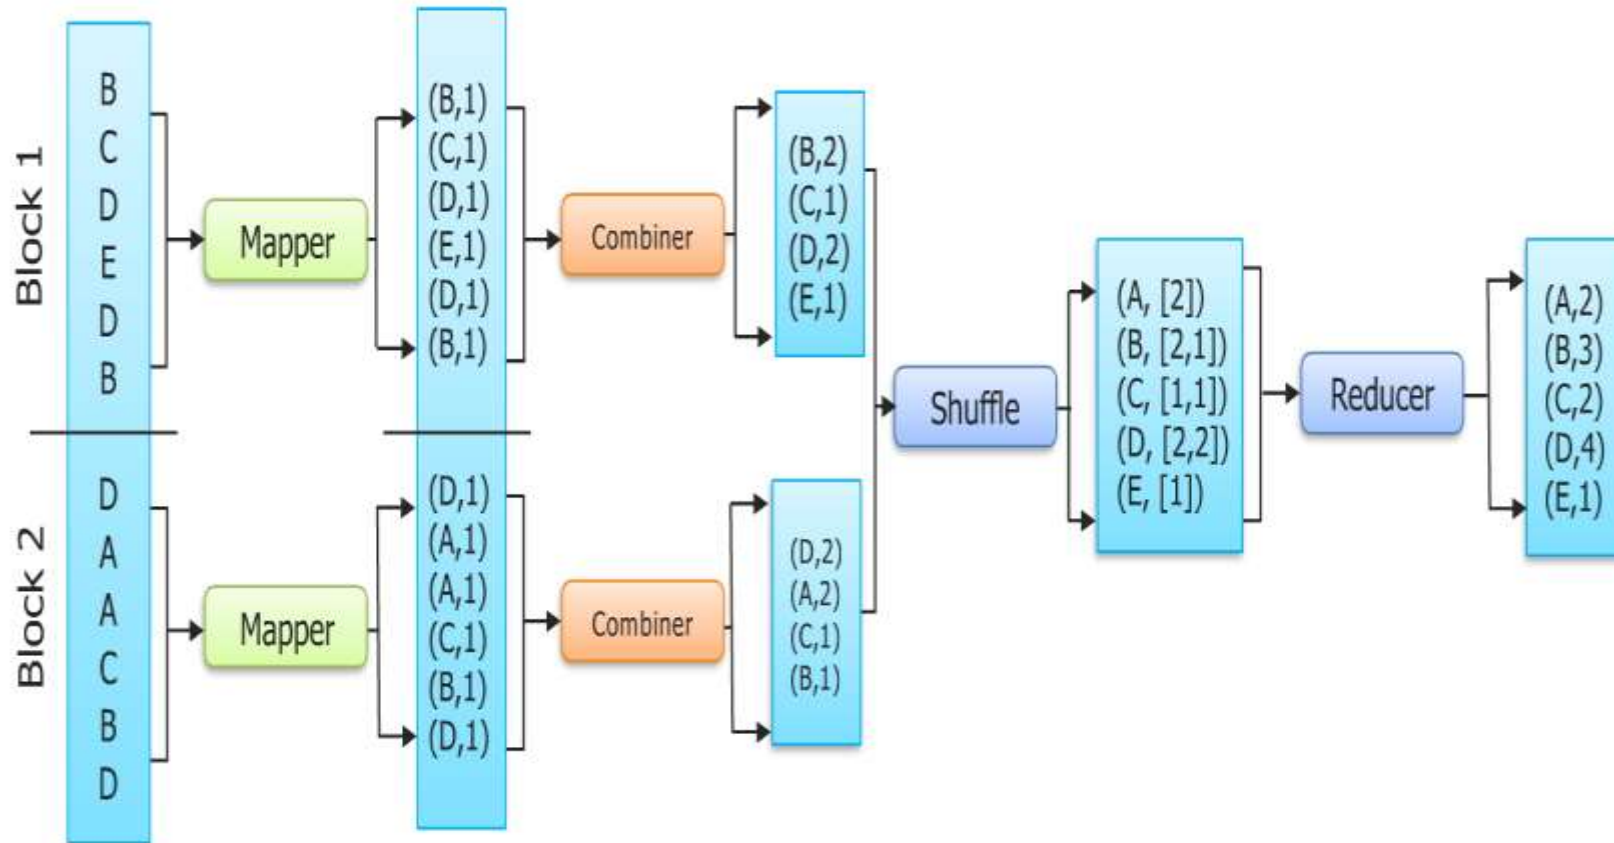
# Overview of map reduce




# Combiner:Local reduce



# Combiner



# Parttitioning:redirecting output from Mapper



## 2. Why should I go for Pig when there is MR?

### → Map-Reduce

- » Powerful model for parallelism.
- » Based on a rigid procedural structure.
- » Provides a good opportunity to parallelize algorithm.



### → PIG

- » It is desirable to have a higher level declarative language.
- » Similar to SQL query where the user specifies the "what" and leaves the "how" to the underlying processing engine.



# What is Pig?

---



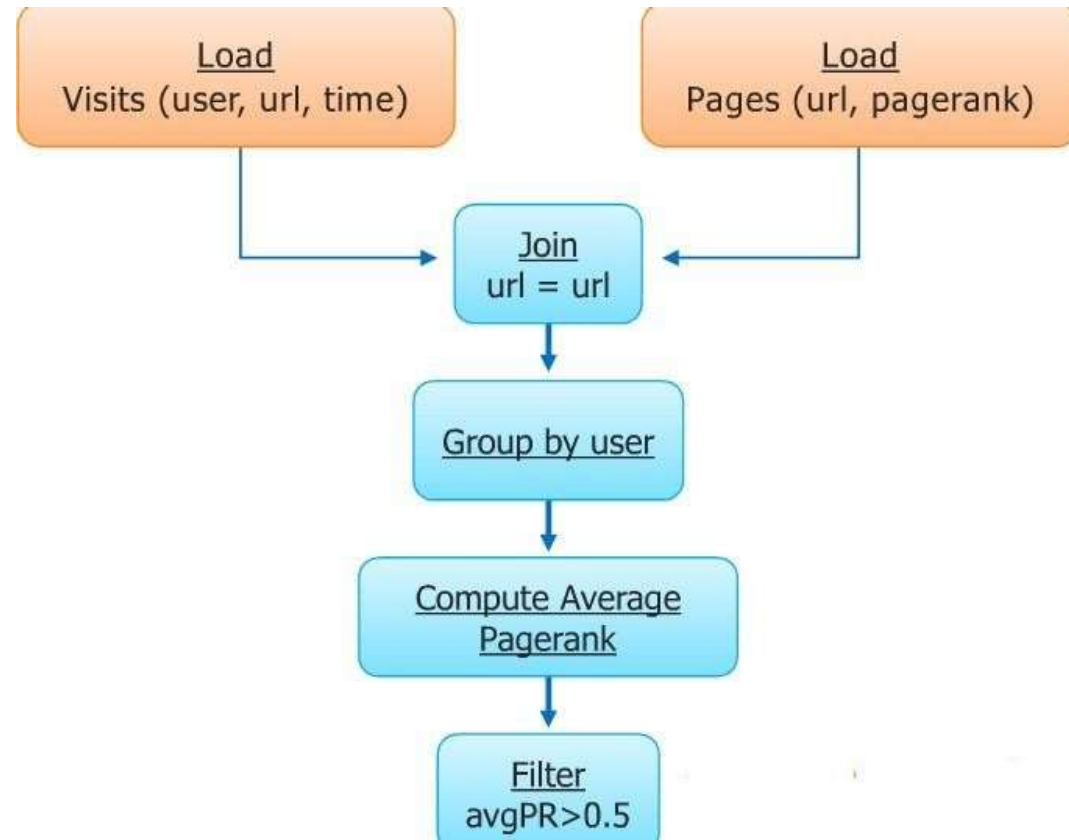
Pig is an [open-source high-level dataflow system](#).

It provides a simple language for queries and data manipulation [Pig Latin](#), that is compiled into map-reduce jobs that are run on Hadoop.

## [Why is it Important?](#)

- Companies like Yahoo, Google and Microsoft are collecting enormous data sets in the form of click streams, search logs, and web crawls.
- Some form of ad-hoc processing and analysis of all of this information is required.

# Conceptual data flow





# Pig basic program structure

## Script:

Pig can run a script file that contains Pig commands.


*Example:* `pig script.pig` runs the commands in the local file `script.pig`.

## Grunt:

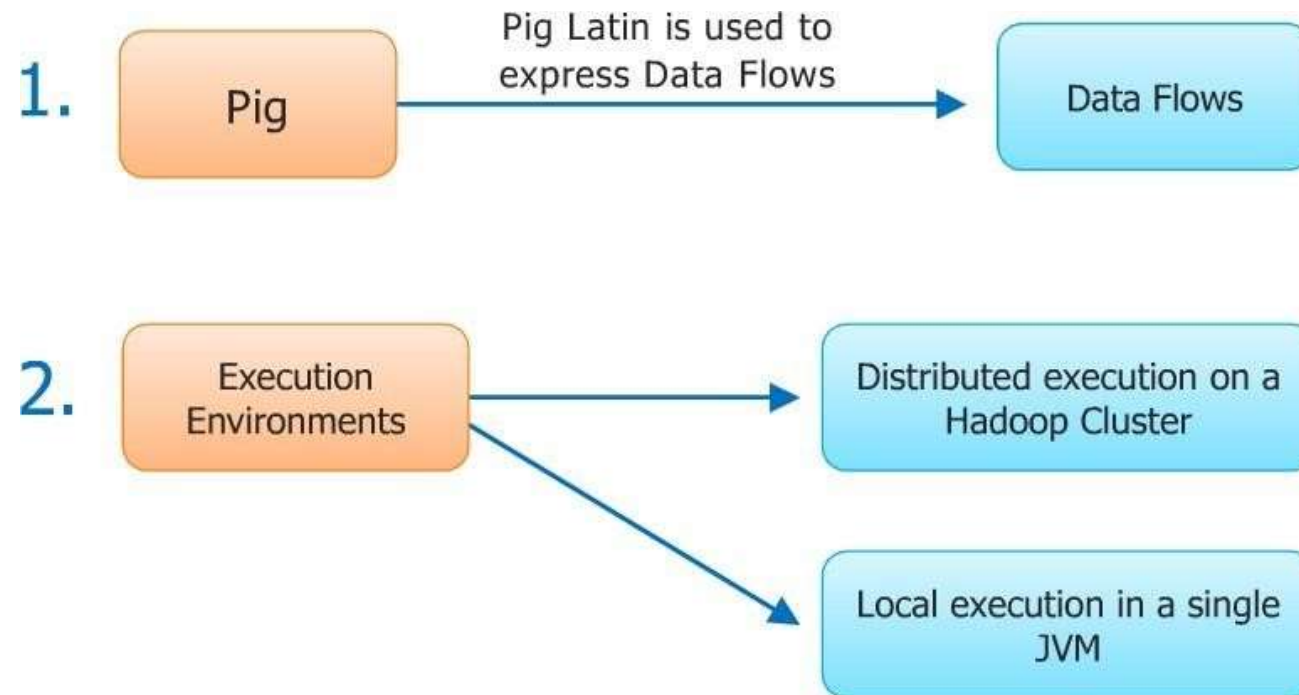
Grunt is an interactive shell for running Pig commands. It is also possible to run Pig scripts from within Grunt using `run` and `exec` (execute).

## Embedded:

Embedded can run Pig programs from Java, much like you can use JDBC to run SQL programs from Java.



# A pig is made of two components




# Pig execution

Pig resides on user machine

Job executes on Cluster




User Machine



Hadoop  
Cluster

No need to install anything extra on your Hadoop Cluster!

# Four basic types of data models



# Data models

Data Models can be defined as follows:

- A **bag** is a collection of tuples.
- A **tuple** is an ordered set of fields.
- A **field** is a piece of data.
- A **Data Map** is a map from keys that are string literals to values that can be any data type.

Example:

```
t = ( 1, {(2,3),(4,6),(5,7)}, ['apache':'search'] )
```

# Pig data types

Pig Data Type	Implementing Class
Bag	org.apache.pig.data.DataBag
Tuple	org.apache.pig.data.Tuple
Map	java.util.Map<Object, Object>
Integer	java.lang.Integer
Long	java.lang.Long
Float	java.lang.Float
Double	java.lang.Double
Chararray	java.lang.String
Bytearray	byte[]

# Pig Latin-Relational operators

Category	Operator	Description
Loading and Storing	LOAD STORE DUMP	Loads data from the file system or other storage into a relation . Saves a relation to the file system or other storage. Prints a relation to the console.
Filtering	FILTER DISTINCT FOREACH...GENERATE STREAM	Removes unwanted rows from a relation. Removes duplicate rows from a relation. Adds or removes fields from a relation. Transforms a relation using an external program.
Grouping and Joining	JOIN COGROUP GROUP CROSS	Joins two or more relations. Groups the data in two or more relations. Groups the data in a single relation. Creates the cross product of two or more relations.
Sorting	ORDER LIMIT	Sorts a relation by one or more fields. Limits the size of a relation to a maximum number of tuples.
Combining and Splitting	UNION SPLIT	Combines two or more relations into one. Splits a relation into two or more relations.

# Pig Latin File Loaders

## Pig Latin File Loaders

`BinStorage` - "binary" storage

`PigStorage` - Loads and stores data that is delimited by something

`TextLoader` - Loads data line by line (delimited by the newline character)

`CSVLoader` - Loads CSV files

`XML Loader` - Loads XML files



# Data

File – Student

Name	Age	GPA
Joe	18	2.5
Sam		3.0
Angel	21	7.9
John	17	9.0
Joe	19	2.9

File – Student Roll

Name	Roll No.
Joe	45
Sam	24
Angel	1
John	12
Joe	19

# Group operator

## Example of GROUP Operator:

```
A = load '/student' USING PigStorage( ',' ) as (name:chararray, age:int, gpa:float);  
dump A;
```

```
(joe,18,2.5)  
(sam,,3.0)  
(angel,21,7.9)  
(john,17,9.0)  
(joe,19,2.9)
```

```
X = group A by name;  
dump X;
```

```
(joe,{(joe,18,2.5),(joe,19,2.9)})  
(sam,{(sam,,3.0)})  
(john,{(john,17,9.0)})  
(angel,{(angel,21,7.9)})
```

# Cogroup operator

Example of COGROU Operator:

```
A = load '/student' USING PigStorage( ',' ) as (name:chararray, age:int,gpa:float);  
B = load '/studentRoll' USING PigStorage( ',' ) as (name:chararray, rollno:int);
```

```
X = cogroup A by name, B by name;  
dump X;
```

```
(joe,{(joe,18,2.5),(joe,19,2.9)},{(joe,45),(joe,19)})  
(sam,{(sam,,3.0)},{(sam,24)})  
(john,{(john,17,9.0)},{(john,12)})  
(angel,{(angel,21,7.9)},{(angel,1)})
```

# Union

**UNION:** To merge the contents of two or more relations.

```
A = LOAD 'data' AS (a1:int,a2:int);

DUMP A;
(1,2)
(4,2)

B = LOAD 'data' AS (b1:int,b2:int);

DUMP B;
(2,4)
(8,9)
(1,3)

X = UNION A, B;

DUMP X;
(2,4)
(8,9)
(1,3)
(1,2)
(4,2)
```

# 3. Real Time Analytics



*Banking*



*Government*



*Healthcare*

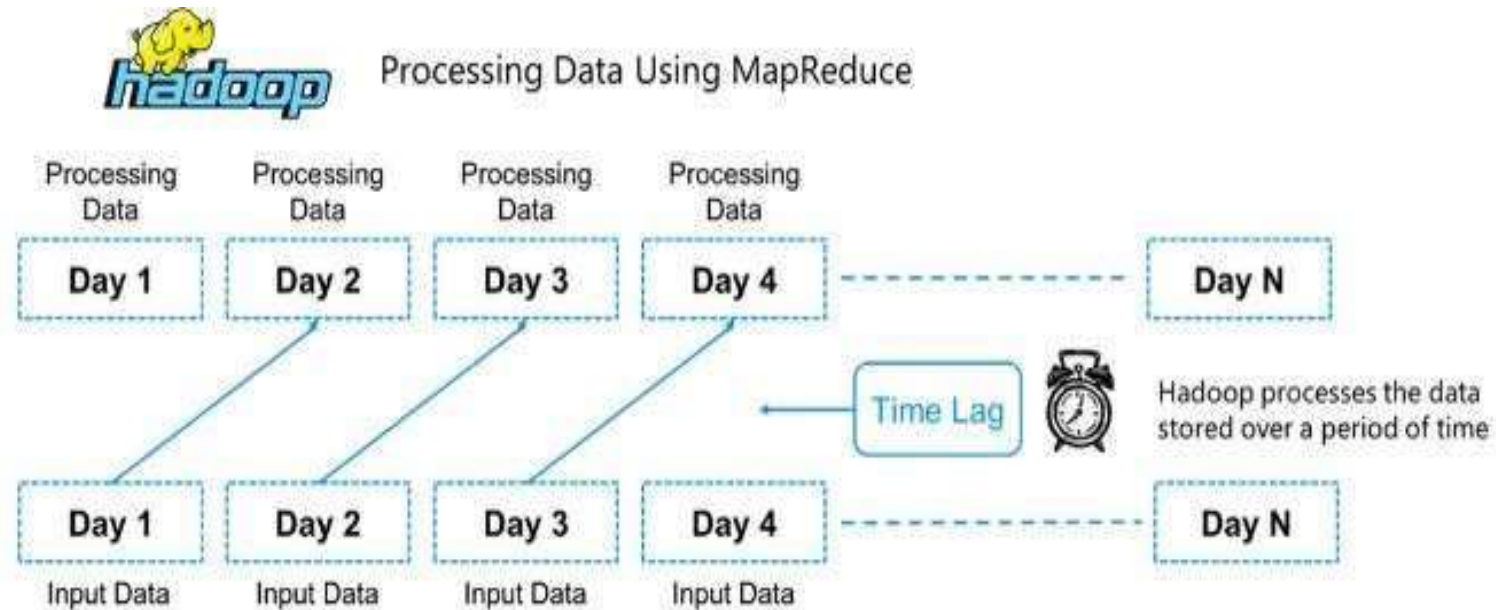


*Telecommunications*

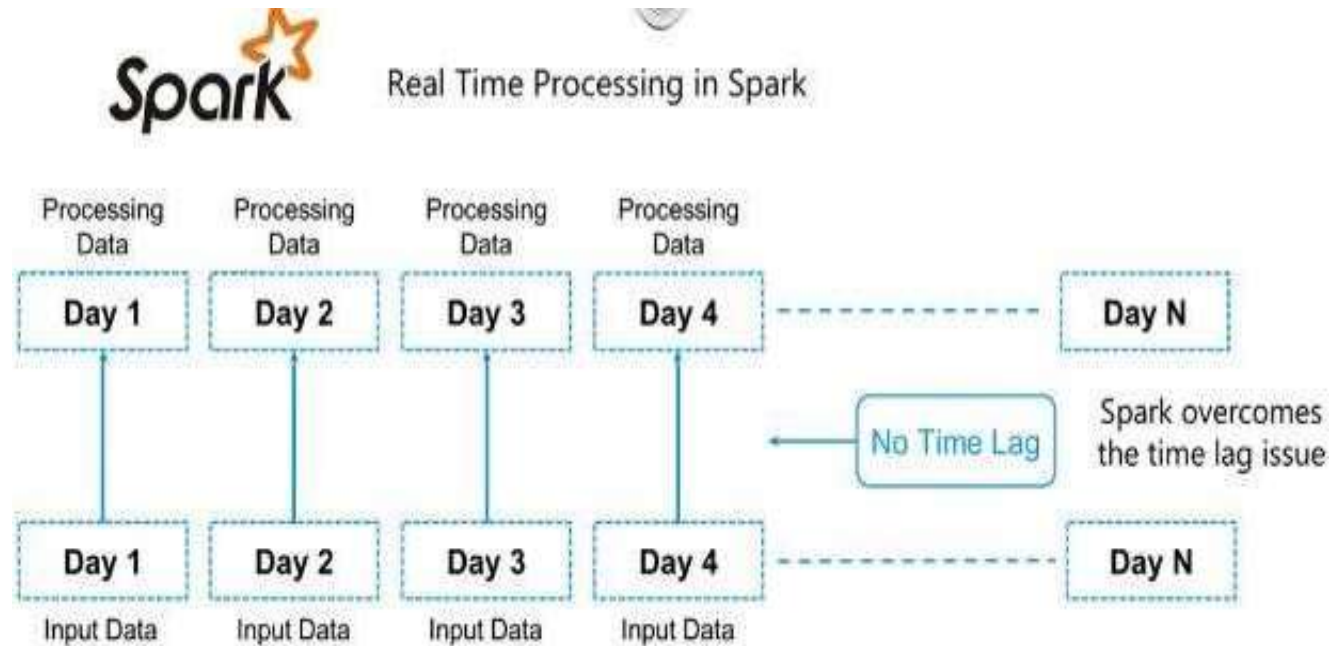


*Stock Market*

# Why Spark when Hadoop is already there?



# Why Spark when Hadoop is already there?



# What is Spark?

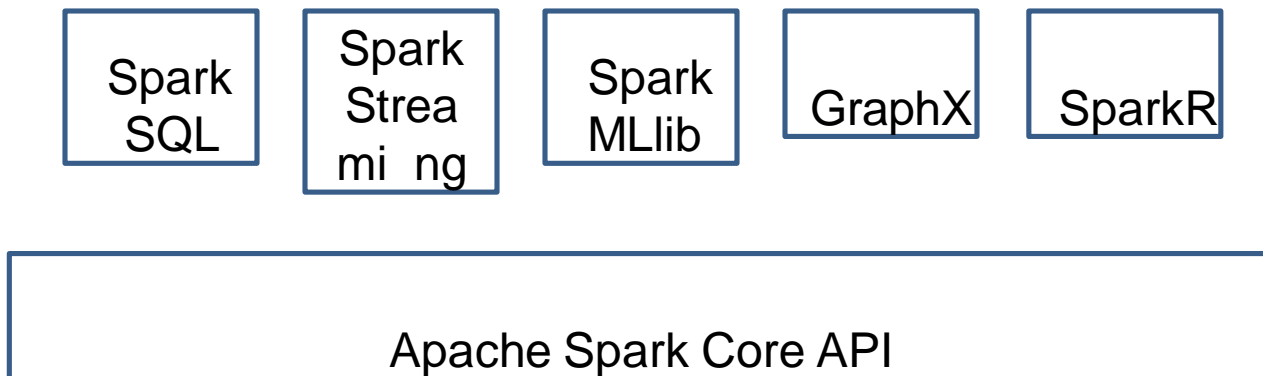
- Apache Spark is an open source cluster computing framework for real-time data processing.
- Main feature of Apache Spark : in-memory cluster computing that increases the processing speed of an application.
- Spark provides an interface for programming entire clusters with implicit data parallelism and fault tolerance.



# Features of Apache Spark



# Spark Components



# Spark Deployment Modes

- Standalone (used for learning & development)
- Local mode (used for learning & development)
- Cluster mode (can work with MESOS or YARN)

# Resilient Distributed Dataset(RDD)

- RDDs are the building blocks of any Spark application. RDDs Stands for:
- **Resilient:** Fault tolerant and is capable of rebuilding data on failure
- **Distributed:** Distributed data among the multiple nodes in a cluster
- **Dataset:** Collection of partitioned data with values

Thank you